

11.2 Health Check and Monitoring Scripts

```
bash
```

```
#!/bin/bash
# scripts/health-check.sh

set -e

echo "🔍 Running comprehensive health checks..."

# Check if all services are running
echo "💻 Checking service status..."
services=("api-gateway" "auth-service" "etsy-service" "design-service" "analytics-service" "postgres" "redis" "")

for service in "${services[@]}"; do
    if docker-compose ps | grep -q "$service.*Up"; then
        echo "✅ $service is running"
    else
        echo "❌ $service is not running"
        exit 1
    fi
done

# Check service endpoints
echo "🌐 Testing service endpoints..."

# API Gateway health check
if curl -f http://localhost:8080/health &>/dev/null; then
    echo "✅ API Gateway health check passed"
else
    echo "❌ API Gateway health check failed"
    exit 1
fi

# Database connectivity
if docker-compose exec -T postgres pg_isready -U postgres &>/dev/null; then
    echo "✅ Database connectivity check passed"
else
    echo "❌ Database connectivity check failed"
    exit 1
fi

# Redis connectivity
if docker-compose exec -T redis redis-cli ping | grep -q "PONG"; then
    echo "✅ Redis connectivity check passed"
else
```

```
echo "✖ Redis connectivity check failed"
exit 1
fi

# MinIO health check
if curl -f http://localhost:9000/minio/health/live &>/dev/null; then
    echo "✓ MinIO health check passed"
else
    echo "✖ MinIO health check failed"
    exit 1
fi

echo "🎉 All health checks passed!"
```

11.3 Automated Backup System

```
python
```

```
# scripts/backup_manager.py
#!/usr/bin/env python3

import os
import sys
import subprocess
import datetime
import boto3
import logging
from pathlib import Path

logging.basicConfig(level=logging.INFO)
logger = logging.getLogger(__name__)

class BackupManager:
    def __init__(self, qnap_path: str = "/share/Container/etsy-saas-prod"):
        self.qnap_path = qnap_path
        self.backup_path = "/share/Container/backups"
        self.timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")

    def backup_database(self):
        """Create PostgreSQL database backup"""
        logger.info("⌚ Starting database backup...")

        backup_file = f"{self.backup_path}/postgres_backup_{self.timestamp}.sql"

        cmd = [
            "docker-compose", "-f", f"{self.qnap_path}/docker-compose.prod.yml",
            "exec", "-T", "postgres",
            "pg_dumpall", "-U", "postgres"
        ]

        try:
            with open(backup_file, 'w') as f:
                subprocess.run(cmd, stdout=f, check=True)

            logger.info(f"✅ Database backup completed: {backup_file}")

            # Compress backup
            subprocess.run(["gzip", backup_file], check=True)
            logger.info(f"✅ Database backup compressed: {backup_file}.gz")

        return f"{backup_file}.gz"
```

```
except subprocess.CalledProcessError as e:
    logger.error(f"❌ Database backup failed: {e}")
    sys.exit(1)

def backup_file_storage(self):
    """Backup MinIO file storage"""
    logger.info("💻 Starting file storage backup...")

    backup_file = f"{self.backup_path}/minio_backup_{self.timestamp}.tar.gz"
    minio_data_path = f"{self.qnap_path}/volumes/minio_data"

    try:
        cmd = ["tar", "-czf", backup_file, "-C", minio_data_path, "."]
        subprocess.run(cmd, check=True)

        logger.info(f"✅ File storage backup completed: {backup_file}")
        return backup_file

    except subprocess.CalledProcessError as e:
        logger.error(f"❌ File storage backup failed: {e}")
        sys.exit(1)

def backup_application_config(self):
    """Backup application configuration and docker-compose files"""
    logger.info("⚙️ Starting configuration backup...")

    backup_file = f"{self.backup_path}/config_backup_{self.timestamp}.tar.gz"

    files_to_backup = [
        "docker-compose.prod.yml",
        "docker-compose.yml",
        ".env.production",
        "traefik",
        "nginx"
    ]

    try:
        cmd = ["tar", "-czf", backup_file, "-C", self.qnap_path]
        cmd.extend(files_to_backup)
        subprocess.run(cmd, check=True)

        logger.info(f"✅ Configuration backup completed: {backup_file}")
        return backup_file

    except subprocess.CalledProcessError as e:
        logger.error(f"❌ Configuration backup failed: {e}")
        sys.exit(1)
```

```
except subprocess.CalledProcessError as e:  
    logger.error(f"❌ Configuration backup failed: {e}")  
    sys.exit(1)  
  
def cleanup_old_backups(self, retention_days: int = 30):  
    """Remove backups older than retention period"""  
    logger.info(f"🧹 Cleaning up backups older than {retention_days} days...")  
  
    cutoff_date = datetime.datetime.now() - datetime.timedelta(days=retention_days)  
  
    for backup_file in Path(self.backup_path).glob("*_backup_*"):  
        if backup_file.stat().st_mtime < cutoff_date.timestamp():  
            backup_file.unlink()  
            logger.info(f"🗂️ Removed old backup: {backup_file}")  
  
def run_full_backup(self):  
    """Run complete backup procedure"""  
    logger.info("🚀 Starting full backup procedure...")  
  
    # Create backup directory  
    os.makedirs(self.backup_path, exist_ok=True)  
  
    # Run all backup procedures  
    db_backup = self.backup_database()  
    file_backup = self.backup_file_storage()  
    config_backup = self.backup_application_config()  
  
    # Cleanup old backups  
    self.cleanup_old_backups()  
  
    logger.info("✅ Full backup completed successfully!")  
  
    return {  
        "database_backup": db_backup,  
        "file_backup": file_backup,  
        "config_backup": config_backup,  
        "timestamp": self.timestamp  
    }  
  
if __name__ == "__main__":  
    backup_manager = BackupManager()  
    backup_manager.run_full_backup()
```

12. Advanced Monitoring and Observability

12.1 Prometheus Metrics Collection

```
python
```

```
# services/common/metrics.py
from prometheus_client import Counter, Histogram, Gauge, start_http_server
import time
from functools import wraps

# Metrics definitions
REQUEST_COUNT = Counter('http_requests_total',
    'Total HTTP requests',
    ['method', 'endpoint', 'status', 'tenant_id'])

REQUEST_DURATION = Histogram('http_request_duration_seconds',
    'HTTP request duration',
    ['method', 'endpoint', 'tenant_id'])

ACTIVE_CONNECTIONS = Gauge('active_connections_total',
    'Active database connections',
    ['tenant_id'])

ETSY_API_CALLS = Counter('etsy_api_calls_total',
    'Total Etsy API calls',
    ['tenant_id', 'endpoint', 'status'])

JOB_PROCESSING_TIME = Histogram('job_processing_seconds',
    'Job processing time',
    ['tenant_id', 'job_type'])

TENANT_RESOURCE_USAGE = Gauge('tenant_resource_usage',
    'Tenant resource usage',
    ['tenant_id', 'resource_type'])

def track_metrics(func):
    """Decorator to track API endpoint metrics"""
    @wraps(func)
    async def wrapper(request, *args, **kwargs):
        start_time = time.time()
        tenant_id = getattr(request.state, 'tenant_id', 'unknown')

        try:
            response = await func(request, *args, **kwargs)
            status = getattr(response, 'status_code', 200)

            # Record metrics
            REQUEST_COUNT.labels(
                method=method,
                endpoint=endpoint,
                status=status,
                tenant_id=tenant_id).inc()
            REQUEST_DURATION.labels(
                method=method,
                endpoint=endpoint,
                tenant_id=tenant_id).observe(time.time() - start_time)
            ACTIVE_CONNECTIONS.labels(tenant_id=tenant_id).inc()
            ETSY_API_CALLS.labels(
                tenant_id=tenant_id,
                endpoint=endpoint,
                status=status).inc()
            JOB_PROCESSING_TIME.labels(
                tenant_id=tenant_id,
                job_type=job_type).observe(time.time() - start_time)
            TENANT_RESOURCE_USAGE.labels(
                tenant_id=tenant_id,
                resource_type=resource_type).set(resource_usage)

        except Exception as e:
            logger.error(f'Error processing request: {e}')
            raise

    return wrapper
```

```
    method=request.method,
    endpoint=request.url.path,
    status=status,
    tenant_id=tenant_id
).inc()

REQUEST_DURATION.labels(
    method=request.method,
    endpoint=request.url.path,
    tenant_id=tenant_id
).observe(time.time() - start_time)

return response

except Exception as e:
    REQUEST_COUNT.labels(
        method=request.method,
        endpoint=request.url.path,
        status=500,
        tenant_id=tenant_id
    ).inc()
    raise

return wrapper

class TenantMetrics:
    def __init__(self, tenant_id: str):
        self.tenant_id = tenant_id

    def record_etsy_api_call(self, endpoint: str, status_code: int):
        """Record Etsy API call metrics"""
        ETSY_API_CALLS.labels(
            tenant_id=self.tenant_id,
            endpoint=endpoint,
            status=status_code
        ).inc()

    def record_job_processing_time(self, job_type: str, duration: float):
        """Record job processing time"""
        JOB_PROCESSING_TIME.labels(
            tenant_id=self.tenant_id,
            job_type=job_type
        ).observe(duration)
```

```
def update_resource_usage(self, resource_type: str, current_usage: int):
    """Update current resource usage"""
    TENANT_RESOURCE_USAGE.labels(
        tenant_id=self.tenant_id,
        resource_type=resource_type
    ).set(current_usage)

# Start metrics server
def start_metrics_server(port: int = 9090):
    """Start Prometheus metrics server"""
    start_http_server(port)
```

12.2 Comprehensive Logging System

python

```
# services/common/logging_config.py

import logging
import json
import sys
from datetime import datetime
from typing import Dict, Any

class TenantFormatter(logging.Formatter):
    """Custom formatter that includes tenant context"""

    def format(self, record):
        # Add tenant context if available
        if hasattr(record, 'tenant_id'):
            record.tenant_context = f"[Tenant: {record.tenant_id}]"
        else:
            record.tenant_context = "[System]"

        # Create structured log entry
        log_entry = {
            "timestamp": datetime.utcnow().isoformat(),
            "level": record.levelname,
            "service": getattr(record, 'service_name', 'unknown'),
            "tenant_id": getattr(record, 'tenant_id', None),
            "message": record.getMessage(),
            "module": record.module,
            "function": record.funcName,
            "line": record.lineno
        }

        # Add extra fields if present
        if hasattr(record, 'extra_fields'):
            log_entry.update(record.extra_fields)

        # Add exception info if present
        if record.exc_info:
            log_entry["exception"] = self.formatException(record.exc_info)

        return json.dumps(log_entry)

class TenantLogger:
    """Logger with tenant context"""

    def __init__(self, service_name: str, tenant_id: str = None):
```

```
self.service_name = service_name
self.tenant_id = tenant_id
self.logger = logging.getLogger(f"{service_name}.{tenant_id or 'system'}")

if not self.logger.handlers:
    self._setup_logger()

def _setup_logger(self):
    """Setup logger with proper formatting"""
    handler = logging.StreamHandler(sys.stdout)
    formatter = TenantFormatter()
    handler.setFormatter(formatter)

    self.logger.addHandler(handler)
    self.logger.setLevel(logging.INFO)

def _log_with_context(self, level: str, message: str, **kwargs):
    """Log message with tenant context"""
    extra = {
        'service_name': self.service_name,
        'tenant_id': self.tenant_id,
        'extra_fields': kwargs
    }

    getattr(self.logger, level.lower())(message, extra=extra)

def info(self, message: str, **kwargs):
    self._log_with_context('INFO', message, **kwargs)

def error(self, message: str, **kwargs):
    self._log_with_context('ERROR', message, **kwargs)

def warning(self, message: str, **kwargs):
    self._log_with_context('WARNING', message, **kwargs)

def debug(self, message: str, **kwargs):
    self._log_with_context('DEBUG', message, **kwargs)

# Usage example in services
def create_tenant_logger(service_name: str):
    def get_logger(tenant_id: str = None):
        return TenantLogger(service_name, tenant_id)
    return get_logger
```

```
# In each service
logger_factory = create_tenant_logger("etsy-service")

@app.middleware("http")
async def logging_middleware(request: Request, call_next):
    tenant_id = getattr(request.state, 'tenant_id', None)
    logger = logger_factory(tenant_id)

    start_time = time.time()

    logger.info("Request started",
               method=request.method,
               path=request.url.path,
               client_ip=request.client.host)

    response = await call_next(request)

    process_time = time.time() - start_time

    logger.info("Request completed",
               method=request.method,
               path=request.url.path,
               status_code=response.status_code,
               process_time=process_time)

    return response
```

13. Security Hardening

13.1 API Security Middleware

```
python
```

```
# services/common/security_middleware.py
from fastapi import Request, HTTPException
from fastapi.security import HTTPBearer, HTTPAuthorizationCredentials
import jwt
import time
import hashlib
from typing import Dict, List
import redis

security = HTTPBearer()
redis_client = redis.Redis.from_url(REDIS_URL)

class SecurityMiddleware:
    def __init__(self):
        self.rate_limits = {
            "api_calls": {"limit": 1000, "window": 3600}, # 1000 calls per hour
            "auth_attempts": {"limit": 5, "window": 900}, # 5 attempts per 15 minutes
        }

    async def rate_limit_check(self, identifier: str, limit_type: str) -> bool:
        """Check if request is within rate limits"""
        config = self.rate_limits.get(limit_type)
        if not config:
            return True

        key = f"rate_limit:{limit_type}:{identifier}"
        current_time = int(time.time())
        window_start = current_time - config["window"]

        # Remove old entries
        redis_client.zremrangebyscore(key, 0, window_start)

        # Count current requests in window
        current_count = redis_client.zcard(key)

        if current_count >= config["limit"]:
            return False

        # Add current request
        redis_client.zadd(key, {str(current_time): current_time})
        redis_client.expire(key, config["window"])

    return True
```

```
async def validate_jwt_token(self, token: str) -> Dict:
    """Validate JWT token and extract claims"""
    try:
        payload = jwt.decode(token, JWT_SECRET, algorithms=["HS256"])

        # Check if token is blacklisted
        token_hash = hashlib.sha256(token.encode()).hexdigest()
        if redis_client.get(f"blacklist:{token_hash}"):
            raise HTTPException(401, "Token has been revoked")

        # Check token expiration
        if payload.get("exp", 0) < time.time():
            raise HTTPException(401, "Token has expired")

    return payload

except jwt.InvalidTokenError:
    raise HTTPException(401, "Invalid token")

async def check_tenant_access(self, token_payload: Dict, requested_tenant: str) -> bool:
    """Verify user has access to requested tenant"""
    token_tenant = token_payload.get("tenant_id")

    if not token_tenant:
        return False

    if token_tenant != requested_tenant:
        # Check if user has multi-tenant access (admin users)
        user_permissions = token_payload.get("permissions", [])
        if "admin:cross_tenant" not in user_permissions:
            return False

    return True

async def log_security_event(self, event_type: str, details: Dict):
    """Log security-related events"""
    security_logger = TenantLogger("security")
    security_logger.warning(f"Security event: {event_type}", **details)

    # Store in Redis for monitoring
    event_key = f"security_events:{int(time.time())}"
    event_data = {
        "type": event_type,
```

```
"timestamp": time.time(),
"details": details
}
redis_client.setex(event_key, 86400, json.dumps(event_data)) # 24 hours

@app.middleware("http")
async def security_middleware(request: Request, call_next):
    security_handler = SecurityMiddleware()

    # Skip security checks for health endpoints
    if request.url.path in ["/health", "/metrics"]:
        return await call_next(request)

    # Rate limiting based on IP
    client_ip = request.client.host
    if not await security_handler.rate_limit_check(client_ip, "api_calls"):
        await security_handler.log_security_event("rate_limit_exceeded", {
            "client_ip": client_ip,
            "path": request.url.path
        })
        raise HTTPException(429, "Rate limit exceeded")

    # Token validation for protected endpoints
    if request.url.path.startswith("/api/"):
        auth_header = request.headers.get("authorization")
        if not auth_header or not auth_header.startswith("Bearer "):
            raise HTTPException(401, "Missing or invalid authorization header")

        token = auth_header.split(" ")[1]
        payload = await security_handler.validate_jwt_token(token)

    # Extract tenant from subdomain or header
    tenant_id = extract_tenant_from_request(request)

    if not await security_handler.check_tenant_access(payload, tenant_id):
        await security_handler.log_security_event("unauthorized_tenant_access", {
            "user_id": payload.get("user_id"),
            "requested_tenant": tenant_id,
            "user_tenant": payload.get("tenant_id"),
            "client_ip": client_ip
        })
        raise HTTPException(403, "Access denied to tenant")

    # Add to request state
```

```
request.state.tenant_id = tenant_id  
request.state.user_id = payload.get("user_id")  
request.state.permissions = payload.get("permissions", [])  
  
response = await call_next(request)  
return response
```

14. Testing Strategy

14.1 Integration Testing Framework

```
python
```

```
# tests/integration/test_multi_tenant.py
import pytest
import asyncio
from httpx import AsyncClient
from sqlalchemy import create_engine
from sqlalchemy.orm import sessionmaker
import docker

@pytest.fixture(scope="session")
def event_loop():
    """Create an instance of the default event loop for the test session."""
    loop = asyncio.get_event_loop_policy().new_event_loop()
    yield loop
    loop.close()

@pytest.fixture(scope="session")
async def test_environment():
    """Set up complete test environment with Docker"""
    client = docker.from_env()

    # Start test containers
    containers = []
    try:
        # Start PostgreSQL
        postgres = client.containers.run(
            "postgres:15",
            environment={
                "POSTGRES_DB": "test_etsy_saas",
                "POSTGRES_USER": "test_user",
                "POSTGRES_PASSWORD": "test_password"
            },
            ports={"5432/tcp": 5433},
            detach=True,
            remove=True
        )
        containers.append(postgres)

    # Start Redis
    redis_container = client.containers.run(
        "redis:7-alpine",
        ports={"6379/tcp": 6380},
        detach=True,
        remove=True
    )

    return {
        "client": client,
        "containers": containers
    }


```

```
)  
    containers.append(redis_container)  
  
    # Wait for services to be ready  
    await asyncio.sleep(10)  
  
    # Set up test database  
    test_engine = create_engine("postgresql://test_user:test_password@localhost:5433/test_etsy_saas")  
  
    yield {  
        "database_url": "postgresql://test_user:test_password@localhost:5433/test_etsy_saas",  
        "redis_url": "redis://localhost:6380",  
        "engine": test_engine  
    }  
  
finally:  
    # Cleanup containers  
    for container in containers:  
        container.stop()  
  
@pytest.fixture  
async def test_tenants(test_environment):  
    """Create test tenants"""  
    engine = test_environment["engine"]  
  
    # Create test schemas and data  
    tenants = [  
        {"id": "test-tenant-1", "subdomain": "testshop1", "name": "Test Shop 1"},  
        {"id": "test-tenant-2", "subdomain": "testshop2", "name": "Test Shop 2"}  
    ]  
  
    with engine.begin() as conn:  
        # Create core tables  
        conn.execute(  
            CREATE TABLE IF NOT EXISTS core.tenants (  
                id UUID PRIMARY KEY,  
                subdomain VARCHAR(63) UNIQUE,  
                company_name VARCHAR(255)  
            )  
        )  
  
        # Insert test tenants  
        for tenant in tenants:  
            conn.execute(  
                INSERT INTO core.tenants (subdomain, name)  
                VALUES (%s, %s)  
            ),  
            (tenant["subdomain"], tenant["name"])  
    ]
```

```
INSERT INTO core.tenants (id, subdomain, company_name)
VALUES (%(id)s, %(subdomain)s, %(name)s)
ON CONFLICT (id) DO NOTHING
"""", tenant)

# Create tenant schema
conn.execute(f"CREATE SCHEMA IF NOT EXISTS tenant_{tenant['subdomain']}")

return tenants

class TestMultiTenantIsolation:
    """Test tenant data isolation"""

    @pytest.mark.asyncio
    async def test_tenant_data_isolation(self, test_environment, test_tenants):
        """Test that tenants cannot access each other's data"""

        async with AsyncClient(base_url="http://localhost:8080") as client:
            # Create data for tenant 1
            tenant1_token = await self.get_tenant_token(client, "test-tenant-1")

            response = await client.post(
                "/api/v1/designs/upload",
                headers={"Authorization": f"Bearer {tenant1_token}"},
                files={"file": ("test.png", b"fake image data", "image/png")}
            )
            assert response.status_code == 201
            design_id = response.json()["id"]

            # Try to access tenant 1's data with tenant 2's token
            tenant2_token = await self.get_tenant_token(client, "test-tenant-2")

            response = await client.get(
                f"/api/v1/designs/{design_id}",
                headers={"Authorization": f"Bearer {tenant2_token}"}
            )
            assert response.status_code == 404 # Should not find the design

    @pytest.mark.asyncio
    async def test_tenant_rate_limiting(self, test_environment, test_tenants):
        """Test that rate limiting is applied per tenant"""

        async with AsyncClient(base_url="http://localhost:8080") as client:
            tenant1_token = await self.get_tenant_token(client, "test-tenant-1")
```

```

tenant2_token = await self.get_tenant_token(client, "test-tenant-2")

# Make requests up to limit for tenant 1
for _ in range(100): # Assume limit is 100 per minute for testing
    response = await client.get(
        "/api/v1/analytics/dashboard",
        headers={"Authorization": f"Bearer {tenant1_token}"}
    )
    if response.status_code == 429:
        break
    else:
        pytest.fail("Rate limit not reached for tenant 1")

# Tenant 2 should still be able to make requests
response = await client.get(
    "/api/v1/analytics/dashboard",
    headers={"Authorization": f"Bearer {tenant2_token}"}
)
assert response.status_code == 200

async def get_tenant_token(self, client: AsyncClient, tenant_id: str) -> str:
    """Helper to get authentication token for tenant"""
    response = await client.post("/api/v1/auth/login", json={
        "email": f"admin@{tenant_id}.com",
        "password": "test_password",
        "tenant_id": tenant_id
    })
    return response.json()["access_token"]

class TestEtsyIntegration:
    """Test Etsy API integration"""

    @pytest.mark.asyncio
    async def test_etsy_oauth_flow(self, test_environment, test_tenants):
        """Test Etsy OAuth integration"""
        # Mock Etsy API responses
        with patch('services.etsy.EtsyService.get_oauth_url') as mock_oauth:
            mock_oauth.return_value = "https://etsy.com/oauth/authorize?..."

        async with AsyncClient(base_url="http://localhost:8080") as client:
            tenant_token = await self.get_tenant_token(client, "test-tenant-1")

            response = await client.get(
                "/api/v1/etsy/oauth-url",

```

```
        headers={"Authorization": f"Bearer {tenant_token}"}
    )

    assert response.status_code == 200
    assert "oauth_url" in response.json()

@pytest.mark.asyncio
async def test_listing_creation(self, test_environment, test_tenants):
    """Test Etsy listing creation"""
    with patch('services.etsy.EtsyService.create_listing') as mock_create:
        mock_create.return_value = {"listing_id": 12345, "status": "active"}

        async with AsyncClient(base_url="http://localhost:8080") as client:
            tenant_token = await self.get_tenant_token(client, "test-tenant-1")

            listing_data = {
                "title": "Test Product",
                "description": "Test description",
                "price": 29.99,
                "quantity": 10
            }

            response = await client.post(
                "/api/v1/etsy/listings",
                headers={"Authorization": f"Bearer {tenant_token}"},
                json=listing_data
            )

            assert response.status_code == 201
            assert response.json()["listing_id"] == 12345

# Performance tests
class TestPerformance:
    """Performance and load testing"""

    @pytest.mark.asyncio
    async def test_concurrent_tenant_requests(self, test_environment, test_tenants):
        """Test handling concurrent requests from multiple tenants"""

        async def make_requests_for_tenant(tenant_id: str, num_requests: int):
            async with AsyncClient(base_url="http://localhost:8080") as client:
                token = await self.get_tenant_token(client, tenant_id)

                tasks = []
```

```

for _ in range(num_requests):
    task = client.get(
        "/api/v1/analytics/dashboard",
        headers={"Authorization": f"Bearer {token}"}
    )
    tasks.append(task)

responses = await asyncio.gather(*tasks)
successful = sum(1 for r in responses if r.status_code == 200)
return successful

# Run concurrent requests for multiple tenants
tasks = [
    make_requests_for_tenant("test-tenant-1", 50),
    make_requests_for_tenant("test-tenant-2", 50)
]

results = await asyncio.gather(*tasks)

# All requests should succeed
assert all(result >= 45 for result in results), "Too many failed requests"

```

15. Migration and Scaling Strategy

15.1 Database Migration System

python

```
# migrations/migration_manager.py
from alembic import command
from alembic.config import Config
from sqlalchemy import create_engine, text
import os

class TenantMigrationManager:
    def __init__(self, database_url: str):
        self.database_url = database_url
        self.engine = create_engine(database_url)

    def run_core_migrations(self):
        """Run migrations for core (non-tenant) tables"""
        alembic_cfg = Config("alembic.ini")
        alembic_cfg.set_main_option("sqlalchemy.url", self.database_url)
        command.upgrade(alembic_cfg, "head")

    def create_tenant_schema(self, tenant_id: str, schema_name: str):
        """Create schema and tables for new tenant"""

        with self.engine.begin() as conn:
            # Create schema
            conn.execute(text(f"CREATE SCHEMA IF NOT EXISTS {schema_name}"))

            # Create tenant-specific tables
            tenant_tables = [
                f"""
                CREATE TABLE {schema_name}.orders (
                    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
                    etsy_receipt_id BIGINT UNIQUE,
                    buyer_email VARCHAR(255),
                    total_amount DECIMAL(10,2),
                    currency VARCHAR(3) DEFAULT 'USD',
                    order_date TIMESTAMP,
                    status VARCHAR(50) DEFAULT 'pending',
                    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
                    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
                )
                """,
                f"""
                CREATE TABLE {schema_name}.products (
                    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
                    etsy_listing_id BIGINT UNIQUE,
                    ...
                )
                """
            ]
```

```
        name VARCHAR(255) NOT NULL,
        description TEXT,
        price DECIMAL(10,2),
        quantity INTEGER DEFAULT 0,
        sku VARCHAR(100),
        tags TEXT[],
        images JSONB,
        created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
        updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
    )
""",
f"""
CREATE TABLE {schema_name}.designs (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name VARCHAR(255) NOT NULL,
    file_path VARCHAR(500),
    file_size BIGINT,
    content_type VARCHAR(100),
    tags TEXT[],
    metadata JSONB,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)
""",
f"""
CREATE TABLE {schema_name}.mockups (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    design_id UUID REFERENCES {schema_name}.designs(id),
    template_name VARCHAR(100),
    mockup_url VARCHAR(500),
    settings JSONB,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)
""",
f"""
CREATE TABLE {schema_name}.jobs (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    job_type VARCHAR(100) NOT NULL,
    status VARCHAR(50) DEFAULT 'queued',
    payload JSONB,
    result JSONB,
    error_message TEXT,
    started_at TIMESTAMP,
    completed_at TIMESTAMP,

```

```
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)
"""

]

for table_sql in tenant_tables:
    conn.execute(text(table_sql))

# Create indexes for performance
indexes = [
    f"CREATE INDEX IF NOT EXISTS idx_{schema_name}_orders_date ON {schema_name}.orders(order_da
f"CREATE INDEX IF NOT EXISTS idx_{schema_name}_orders_status ON {schema_name}.orders(status)
f"CREATE INDEX IF NOT EXISTS idx_{schema_name}_products_etsy_id ON {schema_name}.products(e
f"CREATE INDEX IF NOT EXISTS idx_{schema_name}_designs_created ON {schema_name}.designs(cre
f"CREATE INDEX IF NOT EXISTS idx_{schema_name}_jobs_status ON {schema_name}.jobs(status, crea
]

for index_sql in indexes:
    conn.execute(text(index_sql))

print(f"✓ Created schema and tables for tenant: {tenant_id}")

def migrate_tenant_data(self, from_schema: str, to_schema: str):
    """Migrate data between tenant schemas"""

    tables = ["orders", "products", "designs", "mockups", "jobs"]

    with self.engine.begin() as conn:
        for table in tables:
            # Copy data from old schema to new schema
            copy_sql = f"""
                INSERT INTO {to_schema}.{table}
                SELECT * FROM {from_schema}.{table}
                ON CONFLICT DO NOTHING
            """
            conn.execute(text(copy_sql))

    print(f"✓ Migrated data from {from_schema} to {to_schema}")

def get_tenant_schema_size(self, schema_name: str) -> dict:
    """Get storage usage statistics for tenant schema"""

    with self.engine.connect() as conn:
        result = conn.execute(text(f"""
            SELECT table_name, pg_sizeof_bytea_head(pg_total_relation_size(table_name))
            FROM information_schema.tables
            WHERE table_schema = '{schema_name}'
        """))
        return {row[0]: row[1] for row in result.fetchall()}
```

```

SELECT
    schemaname,
    tablename,
    pg_size.pretty(pg_total_relation_size(schemaname||'.'||tablename)) as size,
    pg_total_relation_size(schemaname||'.'||tablename) as size_bytes
FROM pg_tables
WHERE schemaname = '{schema_name}'
ORDER BY pg_total_relation_size(schemaname||'.'||tablename) DESC
"""))

```

```

tables = []
total_size = 0

```

```

for row in result:
    table_info = {
        "table_name": row.tablename,
        "size_pretty": row.size,
        "size_bytes": row.size_bytes
    }
    tables.append(table_info)
    total_size += row.size_bytes

```

```

return {
    "schema": schema_name,
    "tables": tables,
    "total_size_bytes": total_size,
    "total_size_pretty": self._format_bytes(total_size)
}

```

```

def _format_bytes(self, bytes_size: int) -> str:
    """Format bytes to human readable size"""
    for unit in ['B', 'KB', 'MB', 'GB', 'TB']:
        if bytes_size < 1024.0:
            return f"{bytes_size:.1f} {unit}"
        bytes_size /= 1024.0
    return f"{bytes_size:.1f} PB"

```

```

# Migration CLI tool
if __name__ == "__main__":
    import sys
    import argparse

```

```

parser = argparse.ArgumentParser(description='Tenant Migration Manager')
parser.add_argument('command', choices=['create', 'migrate', 'stats'])

```

```

parser.add_argument('--tenant-id', required=True)
parser.add_argument('--from-schema', help='Source schema for migration')
parser.add_argument('--to-schema', help='Target schema for migration')

args = parser.parse_args()

database_url = os.getenv('DATABASE_URL', 'postgresql://postgres:password@localhost:5432/etsy_saas')
manager = TenantMigrationManager(database_url)

if args.command == 'create':
    schema_name = f"tenant_{args.tenant_id}"
    manager.create_tenant_schema(args.tenant_id, schema_name)

elif args.command == 'migrate':
    if not args.from_schema or not args.to_schema:
        print("Error: --from-schema and --to-schema required for migration")
        sys.exit(1)
    manager.migrate_tenant_data(args.from_schema, args.to_schema)

elif args.command == 'stats':
    schema_name = f"tenant_{args.tenant_id}"
    stats = manager.get_tenant_schema_size(schema_name)

    print(f"\n📊 Storage statistics for {stats['schema']}:")
    print(f"Total size: {stats['total_size_pretty']}")
    print("\nTable breakdown:")
    for table in stats['tables']:
        print(f" {table['table_name']}: {table['size_pretty']}")

```

15.2 Horizontal Scaling Strategy

python

```
# services/scaling/tenant_balancer.py
import asyncio
import aiohttp
from typing import List, Dict, Any
from dataclasses import dataclass
import redis
import json

@dataclass
class ServiceNode:
    id: str
    host: str
    port: int
    capacity: int
    current_load: int
    health_status: str
    tenant_assignments: List[str]

class TenantLoadBalancer:
    """Manages tenant distribution across multiple service nodes"""

    def __init__(self, redis_url: str):
        self.redis = redis.Redis.from_url(redis_url)
        self.nodes: Dict[str, ServiceNode] = {}
        self.tenant_assignments: Dict[str, str] = {} # tenant_id -> node_id

    async def register_node(self, node: ServiceNode):
        """Register a new service node"""
        self.nodes[node.id] = node

        # Store in Redis for persistence
        node_data = {
            "id": node.id,
            "host": node.host,
            "port": node.port,
            "capacity": node.capacity,
            "current_load": node.current_load,
            "health_status": node.health_status,
            "tenant_assignments": node.tenant_assignments
        }

        await self.redis.hset("service_nodes", node.id, json.dumps(node_data))
        print(f"✓ Registered service node: {node.id}")
```

```

async def assign_tenant_to_node(self, tenant_id: str) -> str:
    """Assign tenant to the best available node"""

    # Find node with lowest load that has capacity
    best_node = None
    best_score = float('inf')

    for node in self.nodes.values():
        if node.health_status != "healthy":
            continue

        if len(node.tenant_assignments) >= node.capacity:
            continue

        # Calculate load score (current load + estimated tenant load)
        estimated_tenant_load = await self._estimate_tenant_load(tenant_id)
        load_score = (node.current_load + estimated_tenant_load) / node.capacity

        if load_score < best_score:
            best_score = load_score
            best_node = node

    if not best_node:
        raise Exception("No available nodes with capacity")

    # Assign tenant to node
    best_node.tenant_assignments.append(tenant_id)
    best_node.current_load += await self._estimate_tenant_load(tenant_id)
    self.tenant_assignments[tenant_id] = best_node.id

    # Update Redis
    await self._update_node_in_redis(best_node)
    await self.redis.hset("tenant_assignments", tenant_id, best_node.id)

    print(f"✓ Assigned tenant {tenant_id} to node {best_node.id}")
    return best_node.id

async def get_node_for_tenant(self, tenant_id: str) -> ServiceNode:
    """Get the service node assigned to a tenant"""

    node_id = self.tenant_assignments.get(tenant_id)

    if not node_id:
        # Check Redis for assignment

```

```
node_id = await self.redis.hget("tenant_assignments", tenant_id)
if node_id:
    self.tenant_assignments[tenant_id] = node_id.decode()

if not node_id:
    raise Exception(f"No node assigned to tenant: {tenant_id}")

return self.nodes.get(node_id)

async def rebalance_tenants(self):
    """Rebalance tenant distribution across nodes"""
    print("⌚ Starting tenant rebalancing...")

    # Calculate current load distribution
    total_capacity = sum(node.capacity for node in self.nodes.values())
    total_load = sum(node.current_load for node in self.nodes.values())

    if total_load == 0:
        return

    target_load_per_node = total_load / len(self.nodes)

    # Identify overloaded and underloaded nodes
    overloaded_nodes = [
        node for node in self.nodes.values()
        if node.current_load > target_load_per_node * 1.2
    ]

    underloaded_nodes = [
        node for node in self.nodes.values()
        if node.current_load < target_load_per_node * 0.8
    ]

    # Move tenants from overloaded to underloaded nodes
    for overloaded_node in overloaded_nodes:
        excess_load = overloaded_node.current_load - target_load_per_node

        # Sort tenants by load (move smallest first for minimal disruption)
        tenant_loads = []
        for tenant_id in overloaded_node.tenant_assignments:
            load = await self._estimate_tenant_load(tenant_id)
            tenant_loads.append((tenant_id, load))

        tenant_loads.sort(key=lambda x: x[1])
```

```

for tenant_id, tenant_load in tenant_loads:
    if excess_load <= 0:
        break

    # Find suitable underloaded node
    target_node = None
    for underloaded_node in underloaded_nodes:
        if (underloaded_node.current_load + tenant_load) <= underloaded_node.capacity:
            target_node = underloaded_node
            break

    if target_node:
        # Migrate tenant
        await self._migrate_tenant(tenant_id, overloaded_node.id, target_node.id)
        excess_load -= tenant_load

print("✅ Tenant rebalancing completed")

async def _migrate_tenant(self, tenant_id: str, from_node_id: str, to_node_id: str):
    """Migrate tenant from one node to another"""
    print(f"🕒 Migrating tenant {tenant_id} from {from_node_id} to {to_node_id}")

    from_node = self.nodes[from_node_id]
    to_node = self.nodes[to_node_id]
    tenant_load = await self._estimate_tenant_load(tenant_id)

    # Update node assignments
    from_node.tenant_assignments.remove(tenant_id)
    from_node.current_load -= tenant_load

    to_node.tenant_assignments.append(tenant_id)
    to_node.current_load += tenant_load

    self.tenant_assignments[tenant_id] = to_node_id

    # Update Redis
    await self._update_node_in_redis(from_node)
    await self._update_node_in_redis(to_node)
    await self.redis.hset("tenant_assignments", tenant_id, to_node_id)

    # Notify application of tenant migration
    await self._notify_tenant_migration(tenant_id, to_node)

```

```

async def _estimate_tenant_load(self, tenant_id: str) -> int:
    """Estimate resource load for a tenant"""
    # Get historical metrics from Redis
    metrics_key = f"tenant_metrics:{tenant_id}"
    metrics_data = await self.redis.get(metrics_key)

    if not metrics_data:
        return 10 # Default load estimate

    metrics = json.loads(metrics_data)

    # Calculate load based on various factors
    api_calls_per_hour = metrics.get("api_calls_per_hour", 0)
    active_jobs = metrics.get("active_jobs", 0)
    data_size_mb = metrics.get("data_size_mb", 0)

    # Load calculation formula (can be refined based on actual usage patterns)
    estimated_load = (
        (api_calls_per_hour / 100) + # Normalize API calls
        (active_jobs * 2) +         # Jobs have higher resource usage
        (data_size_mb / 1000)       # Storage factor
    )

    return max(1, int(estimated_load))

async def _update_node_in_redis(self, node: ServiceNode):
    """Update node information in Redis"""
    node_data = {
        "id": node.id,
        "host": node.host,
        "port": node.port,
        "capacity": node.capacity,
        "current_load": node.current_load,
        "health_status": node.health_status,
        "tenant_assignments": node.tenant_assignments
    }
    await self.redis.hset("service_nodes", node.id, json.dumps(node_data))

async def _notify_tenant_migration(self, tenant_id: str, target_node: ServiceNode):
    """Notify relevant services about tenant migration"""
    notification = {
        "type": "tenant_migration",
        "tenant_id": tenant_id,
        "new_node": {

```

```

        "id": target_node.id,
        "host": target_node.host,
        "port": target_node.port
    },
    "timestamp": asyncio.get_event_loop().time()
}

# Publish to Redis for service consumption
await self.redis.publish("tenant_migrations", json.dumps(notification))

# Auto-scaling based on metrics
class AutoScaler:
    """Automatically scale services based on load metrics"""

    def __init__(self, load_balancer: TenantLoadBalancer):
        self.load_balancer = load_balancer
        self.scaling_cooldown = 300 # 5 minutes between scaling operations
        self.last_scale_time = {}

    async def check_scaling_needs(self):
        """Check if scaling is needed based on current metrics"""
        current_time = asyncio.get_event_loop().time()

        # Calculate overall system load
        total_capacity = sum(
            node.capacity for node in self.load_balancer.nodes.values()
            if node.health_status == "healthy"
        )
        total_load = sum(
            node.current_load for node in self.load_balancer.nodes.values()
            if node.health_status == "healthy"
        )

        if total_capacity == 0:
            return

        utilization = total_load / total_capacity

        # Scale up if utilization > 80%
        if utilization > 0.8:
            if self._can_scale("up", current_time):
                await self._scale_up()
                self.last_scale_time["up"] = current_time

```

```

# Scale down if utilization < 20% and we have more than 1 node
elif utilization < 0.2 and len(self.load_balancer.nodes) > 1:
    if self._can_scale("down", current_time):
        await self._scale_down()
        self.last_scale_time["down"] = current_time

def _can_scale(self, direction: str, current_time: float) -> bool:
    """Check if scaling operation can be performed (respects cooldown)"""
    last_scale = self.last_scale_time.get(direction, 0)
    return (current_time - last_scale) > self.scaling_cooldown

async def _scale_up(self):
    """Add a new service node"""
    print("↗ Scaling up: Adding new service node...")

    # In a real implementation, this would:
    # 1. Launch new container/VM
    # 2. Wait for health check to pass
    # 3. Register with load balancer

    # For demonstration, we'll simulate adding a node
    new_node_id = f"node_{len(self.load_balancer.nodes) + 1}"
    new_node = ServiceNode(
        id=new_node_id,
        host="localhost",
        port=8000 + len(self.load_balancer.nodes),
        capacity=50,
        current_load=0,
        health_status="healthy",
        tenant_assignments=[]
    )

    await self.load_balancer.register_node(new_node)
    print(f"✓ Added new service node: {new_node_id}")

async def _scale_down(self):
    """Remove a service node with minimal disruption"""
    print("↖ Scaling down: Removing service node...")

    # Find node with lowest load for removal
    candidate_node = min(
        self.load_balancer.nodes.values(),
        key=lambda n: n.current_load
    )

```

```

if candidate_node.tenant_assignments:
    # Migrate tenants to other nodes
    for tenant_id in candidate_node.tenant_assignments[:]:
        # Find alternative node
        target_node = min(
            [n for n in self.load_balancer.nodes.values()
             if n.id != candidate_node.id and n.health_status == "healthy"],
            key=lambda n: n.current_load
        )

        await self.load_balancer._migrate_tenant(
            tenant_id, candidate_node.id, target_node.id
        )

    # Remove node
    del self.load_balancer.nodes[candidate_node.id]
    await self.load_balancer.redis.hdel("service_nodes", candidate_node.id)

    print(f"✓ Removed service node: {candidate_node.id}")

# Usage example
async def main():
    load_balancer = TenantLoadBalancer("redis://localhost:6379")
    auto_scaler = AutoScaler(load_balancer)

    # Register initial nodes
    for i in range(2):
        node = ServiceNode(
            id=f"node_{i+1}",
            host="localhost",
            port=8001 + i,
            capacity=50,
            current_load=0,
            health_status="healthy",
            tenant_assignments=[]
        )
        await load_balancer.register_node(node)

    # Assign some tenants
    tenants = ["tenant1", "tenant2", "tenant3", "tenant4"]
    for tenant in tenants:
        await load_balancer.assign_tenant_to_node(tenant)

```

```
# Run auto-scaling check
while True:
    await auto_scaler.check_scaling_needs()
    await load_balancer.rebalance_tenants()
    await asyncio.sleep(60) # Check every minute

if __name__ == "__main__":
    asyncio.run(main())
```

16. Production Deployment Guide

16.1 Complete Production Docker Compose

yaml

```
# docker-compose.production.yml
version: '3.8'

services:
  # Load Balancer / Reverse Proxy
  traefik:
    image: traefik:v2.10
    container_name: etsy-traefik
    command:
      - --api.dashboard=true
      - --api.insecure=true
      - --providers.docker=true
      - --providers.docker.exposedbydefault=false
      - --entrypoints.web.address=:80
      - --entrypoints.websecure.address=:443
      - --certificatesresolvers.letsencrypt.acme.tlschallenge=true
      - --certificatesresolvers.letsencrypt.acme.email=admin@yourdomain.com
      - --certificatesresolvers.letsencrypt.acme.storage=/letsencrypt/acme.json
      - --metrics.prometheus=true
    ports:
      - "80:80"
      - "443:443"
      - "8080:8080" # Dashboard
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock:ro"
      - "./letsencrypt:/letsencrypt"
    networks:
      - etsy-network
    restart: unless-stopped
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U postgres"]
      interval: 30s
      timeout: 5s
      retries: 5

  redis:
    image: redis:7-alpine
    container_name: etsy-redis
    command: redis-server --appendonly yes --maxmemory 1gb --maxmemory-policy allkeys-lru
    volumes:
      - redis_data:/data
    ports:
      - "6379:6379"
```

```
deploy:
resources:
limits:
  memory: 1G
  cpus: '0.5'
reservations:
  memory: 512M
  cpus: '0.25'
networks:
- etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "redis-cli", "ping"]
  interval: 30s
  timeout: 5s
  retries: 3

minio:
image: minio/minio:latest
container_name: etsy-minio
command: server /data --console-address ":9001"
environment:
- MINIO_ROOT_USER=${MINIO_ACCESS_KEY}
- MINIO_ROOT_PASSWORD=${MINIO_SECRET_KEY}
- MINIO_BROWSER_REDIRECT_URL=https://minio.yourdomain.com
volumes:
- minio_data:/data
ports:
- "9000:9000"
- "9001:9001"
deploy:
resources:
limits:
  memory: 1G
  cpus: '1.0'
reservations:
  memory: 512M
  cpus: '0.5'
networks:
- etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:9000/minio/health/live"]
  interval: 30s
```

```
timeout: 5s
retries: 3
labels:
- "traefik.enable=true"
- "traefik.http.routers.minio.rule=Host(`minio.yourdomain.com`)"
- "traefik.http.routers.minio.tls=true"
- "traefik.http.routers.minio.tls.certresolver=letsencrypt"
- "traefik.http.services.minio.loadbalancer.server.port=9001"

# Monitoring Stack
prometheus:
image: prom/prometheus:latest
container_name: etsy-prometheus
command:
- '--config.file=/etc/prometheus/prometheus.yml'
- '--storage.tsdb.path=/prometheus'
- '--web.console.libraries=/etc/prometheus/console_libraries'
- '--web.console.templates=/etc/prometheus/consoles'
- '--storage.tsdb.retention.time=30d'
- '--web.enable-lifecycle'
volumes:
- ./monitoring/prometheus.yml:/etc/prometheus/prometheus.yml
- prometheus_data:/prometheus
ports:
- "9090:9090"
deploy:
resources:
limits:
  memory: 1G
  cpus: '1.0'
reservations:
  memory: 512M
  cpus: '0.5'
networks:
- etsy-network
restart: unless-stopped
labels:
- "traefik.enable=true"
- "traefik.http.routers.prometheus.rule=Host(`prometheus.yourdomain.com`)"
- "traefik.http.routers.prometheus.tls=true"
- "traefik.http.routers.prometheus.tls.certresolver=letsencrypt"

grafana:
image: grafana/grafana:latest
```

```
container_name: etsy-grafana
environment:
  - GF_SECURITY_ADMIN_USER=${GRAFANA_ADMIN_USER}
  - GF_SECURITY_ADMIN_PASSWORD=${GRAFANA_ADMIN_PASSWORD}
  - GF_USERS_ALLOW_SIGN_UP=false
volumes:
  - grafana_data:/var/lib/grafana
  - ./monitoring/grafana/provisioning:/etc/grafana/provisioning
  - ./monitoring/grafana/dashboards:/var/lib/grafana/dashboards
ports:
  - "3000:3000"
deploy:
resources:
limits:
  memory: 512M
  cpus: '0.5'
reservations:
  memory: 256M
  cpus: '0.25'
depends_on:
  - prometheus
networks:
  - etsy-network
restart: unless-stopped
labels:
  - "traefik.enable=true"
  - "traefik.http.routers.grafana.rule=Host(`monitoring.yourdomain.com`)"
  - "traefik.http.routers.grafana.tls=true"
  - "traefik.http.routers.grafana.tls.certresolver=letsencrypt"
```

Log Management

```
loki:
image: grafana/loki:latest
container_name: etsy-loki
command: -config.file=/etc/loki/local-config.yaml
volumes:
  - ./monitoring/loki/config.yaml:/etc/loki/local-config.yaml
  - loki_data:/loki
ports:
  - "3100:3100"
deploy:
resources:
limits:
  memory: 512M
```

```
    cpus: '0.5'
  reservations:
    memory: 256M
  cpus: '0.25'
networks:
- etsy-network
restart: unless-stopped

promtail:
image: grafana/promtail:latest
container_name: etsy-promtail
command: -config.file=/etc/promtail/config.yml
volumes:
- ./monitoring/promtail/config.yml:/etc/promtail/config.yml
- /var/log:/var/log:ro
- /var/run/docker.sock:/var/run/docker.sock:ro
deploy:
resources:
limits:
  memory: 256M
  cpus: '0.25'
reservations:
  memory: 128M
  cpus: '0.1'
depends_on:
- loki
networks:
- etsy-network
restart: unless-stopped

volumes:
postgres_data:
driver: local
driver_opts:
type: none
device: /share/Container/volumes/postgres
o: bind
redis_data:
driver: local
driver_opts:
type: none
device: /share/Container/volumes/redis
o: bind
minio_data:
```

```
driver: local
driver_opts:
  type: none
  device: /share/Container/volumes/minio
  o: bind

prometheus_data:
  driver: local
  driver_opts:
    type: none
    device: /share/Container/volumes/prometheus
    o: bind

grafana_data:
  driver: local
  driver_opts:
    type: none
    device: /share/Container/volumes/grafana
    o: bind

loki_data:
  driver: local
  driver_opts:
    type: none
    device: /share/Container/volumes/loki
    o: bind

networks:
  etsy-network:
    driver: bridge
    ipam:
      config:
        - subnet: 172.20.0.0/16
```

16.2 Environment Configuration

```
bash
```

```
# .env.production
# Database Configuration
POSTGRES_PASSWORD=your_secure_database_password_here
DATABASE_URL=postgresql://postgres:your_secure_database_password_here@postgres:5432/etsy_saas

# Redis Configuration
REDIS_URL=redis://redis:6379

# JWT Security
JWT_SECRET=your_very_secure_jwt_secret_key_minimum_32_chars

# MinIO S3-Compatible Storage
MINIO_ACCESS_KEY=your_minio_access_key
MINIO_SECRET_KEY=your_minio_secret_key_minimum_8_chars

# Monitoring
GRAFANA_ADMIN_USER=admin
GRAFANA_ADMIN_PASSWORD=your_secure_grafana_password

# Etsy API Configuration
ETSY_CLIENT_ID=your_etsy_client_id
ETSY_CLIENT_SECRET=your_etsy_client_secret
ETSY_REDIRECT_URI=https://yourdomain.com/auth/etsy/callback

# Email Configuration (for notifications)
SMTP_HOST=smtp.gmail.com
SMTP_PORT=587
SMTP_USER=your-email@gmail.com
SMTP_PASSWORD=your_app_password
SMTP_FROM_EMAIL=noreply@yourdomain.com

# Stripe Configuration (for billing)
STRIPE_PUBLIC_KEY=pk_live_your_stripe_public_key
STRIPE_SECRET_KEY=sk_live_your_stripe_secret_key
STRIPE_WEBHOOK_SECRET=whsec_your_webhook_secret

# Domain Configuration
PRIMARY_DOMAIN=yourdomain.com
TENANT_SUBDOMAIN_SUFFIX=.yourdomain.com

# Backup Configuration
BACKUP_S3_BUCKET=etsy-saas-backups
AWS_ACCESS_KEY_ID=your_aws_access_key
```

```
AWS_SECRET_ACCESS_KEY=your_aws_secret_key  
AWS_REGION=us-east-1
```

```
# Rate Limiting  
RATE_LIMIT_REQUESTS_PER_MINUTE=1000  
RATE_LIMIT_BURST=100
```

```
# Resource Limits  
MAX_UPLOAD_SIZE_MB=100  
MAX_DESIGNS_PER_TENANT=10000  
MAX_LISTINGS_PER_BATCH=100
```

```
# Feature Flags  
ENABLE_AUTO_SCALING=true  
ENABLE_METRICS_COLLECTION=true  
ENABLE_ADVANCED_ANALYTICS=true  
ENABLE_MULTI_REGION=false
```

16.3 Production Startup Script

```
bash
```

```
#!/bin/bash
# start-production.sh

set -euo pipefail

# Configuration
PROJECT_NAME="etsy-saas"
COMPOSE_FILE="docker-compose.production.yml"
ENV_FILE=".env.production"

# Colors for output
RED='\033[0;31m'
GREEN='\033[0;32m'
YELLOW='\033[1;33m'
BLUE='\033[0;34m'
NC='\033[0m' # No Color

# Logging function
log() {
    echo -e "${GREEN}[$(date +'%Y-%m-%d %H:%M:%S')}${NC} $1"
}

error() {
    echo -e "${RED}[ERROR]${NC} $1" >&2
}

warning() {
    echo -e "${YELLOW}[WARNING]${NC} $1"
}

info() {
    echo -e "${BLUE}[INFO]${NC} $1"
}

# Check if running on QNAP
check_qnap_environment() {
    if [[ ! -d "/share/Container" ]]; then
        error "This script is designed for QNAP Container Station environment"
        exit 1
    fi

    log "✅ QNAP Container Station environment detected"
}
```

```

# Pre-flight checks
pre_flight_checks() {
    log "🔍 Running pre-flight checks..."

    # Check if docker-compose is available
    if ! command -v docker-compose &> /dev/null; then
        error "docker-compose is not installed or not in PATH"
        exit 1
    fi

    # Check if required files exist
    if [[ ! -f "$COMPOSE_FILE" ]]; then
        error "Docker Compose file not found: $COMPOSE_FILE"
        exit 1
    fi

    if [[ ! -f "$ENV_FILE" ]]; then
        error "Environment file not found: $ENV_FILE"
        exit 1
    fi

    # Check available disk space (minimum 10GB)
    available_space=$(df /share/Container | awk 'NR==2 {print $4}')
    required_space=$((10 * 1024 * 1024)) # 10GB in KB

    if [[ $available_space -lt $required_space ]]; then
        error "Insufficient disk space. Required: 10GB, Available: ${((available_space / 1024 / 1024))}GB"
        exit 1
    fi

    # Check available memory (minimum 4GB)
    available_memory=$(free -m | awk 'NR==2 {print $7}')
    required_memory=4096

    if [[ $available_memory -lt $required_memory ]]; then
        warning "Low available memory. Available: ${available_memory}MB, Recommended: ${required_memory}MB"
    fi

    log "✅ Pre-flight checks completed"
}

# Create necessary directories
setup_directories() {

```

```
log "📅 Setting up directory structure..."
```

```
directories=(  
    "/share/Container/volumes/postgres"  
    "/share/Container/volumes/redis"  
    "/share/Container/volumes/minio"  
    "/share/Container/volumes/prometheus"  
    "/share/Container/volumes/grafana"  
    "/share/Container/volumes/loki"  
    "/share/Container/backups"  
    "/share/Container/logs"  
    "./letsencrypt"  
    "./monitoring/prometheus"  
    "./monitoring/grafana/provisioning/dashboards"  
    "./monitoring/grafana/provisioning/datasources"  
    "./monitoring/loki"  
    "./monitoring/promtail"  
)
```

```
for dir in "${directories[@]}"; do  
    mkdir -p "$dir"  
    info "Created directory: $dir"  
done
```

```
# Set proper permissions  
chmod -R 755 /share/Container/volumes  
chmod 600 ./letsencrypt
```

```
log "✅ Directory structure setup completed"
```

```
}
```

```
# Generate monitoring configurations  
setup_monitoring_config() {  
    log "📊 Setting up monitoring configurations..."  
  
    # Prometheus configuration  
    cat > ./monitoring/prometheus/prometheus.yml << EOF  
global:  
    scrape_interval: 15s  
    evaluation_interval: 15s  
  
rule_files:  
    # - "first_rules.yml"
```

```
scrape_configs:
```

```
- job_name: 'prometheus'
```

```
  static_configs:
```

```
    - targets: ['localhost:9090']
```

```
- job_name: 'etsy-services'
```

```
  static_configs:
```

```
    - targets: [
```

```
      'api-gateway:9090',
```

```
      'auth-service:9090',
```

```
      'etsy-service:9090',
```

```
      'design-service:9090',
```

```
      'analytics-service:9090'
```

```
    ]
```

```
scrape_interval: 30s
```

```
metrics_path: '/metrics'
```

```
- job_name: 'traefik'
```

```
  static_configs:
```

```
    - targets: ['traefik:8080']
```

```
- job_name: 'postgres'
```

```
  static_configs:
```

```
    - targets: ['postgres:9187']
```

```
- job_name: 'redis'
```

```
  static_configs:
```

```
    - targets: ['redis:9121']
```

```
EOF
```

```
# Grafana datasource configuration
```

```
mkdir -p ./monitoring/grafana/provisioning/datasources
```

```
cat > ./monitoring/grafana/provisioning/datasources/prometheus.yml << EOF
```

```
apiVersion: 1
```

```
datasources:
```

```
- name: Prometheus
```

```
  type: prometheus
```

```
  access: proxy
```

```
  url: http://prometheus:9090
```

```
  isDefault: true
```

```
  editable: true
```

```
- name: Loki
```

```
type: loki
access: proxy
url: http://loki:3100
editable: true
EOF

# Loki configuration
cat > ./monitoring/loki/config.yaml << EOF
auth_enabled: false

server:
  http_listen_port: 3100
  grpc_listen_port: 9096

common:
  path_prefix: /loki
storage:
  filesystem:
    chunks_directory: /loki/chunks
    rules_directory: /loki/rules
replication_factor: 1
ring:
  instance_addr: 127.0.0.1
  kvstore:
    store: inmemory

query_range:
results_cache:
  cache:
    embedded_cache:
      enabled: true
      max_size_mb: 100

schema_config:
configs:
  - from: 2020-10-24
    store: boltdb-shipper
    object_store: filesystem
    schema: v11
    index:
      prefix: index_
      period: 24h

ruler:
```

```
alertmanager_url: http://localhost:9093
```

```
limits_config:  
  reject_old_samples: true  
  reject_old_samples_max_age: 168h
```

```
chunk_store_config:  
  max_look_back_period: 0s
```

```
table_manager:  
  retention_deletes_enabled: false  
  retention_period: 0s
```

```
compactor:  
  working_directory: /loki/boltdb-shipper-compactor  
  shared_store: filesystem  
  compaction_interval: 10m  
  retention_enabled: true  
  retention_delete_delay: 2h  
  retention_delete_worker_count: 150
```

```
ingester:  
  max_chunk_age: 1h  
  chunk_idle_period: 30s  
  chunk_block_size: 262144  
  chunk_target_size: 1048576  
  chunk_retain_period: 30s  
  max_transfer_retries: 0
```

```
wal:
```

```
  enabled: true  
  dir: /loki/wal
```

```
EOF
```

```
# Promtail configuration  
cat > ./monitoring/promtail/config.yml << EOF
```

```
server:  
  http_listen_port: 9080  
  grpc_listen_port: 0
```

```
positions:  
  filename: /tmp/positions.yaml
```

```
clients:  
  - url: http://loki:3100/loki/api/v1/push
```

```
scrape_configs:
  - job_name: containers
    static_configs:
      - targets:
          - localhost
        labels:
          job: containerlogs
          __path__: /var/log/containers/*.log

  - job_name: docker
    docker_sd_configs:
      - host: unix:///var/run/docker.sock
        refresh_interval: 5s
    relabel_configs:
      - source_labels: ['__meta_docker_container_name']
        regex: '/(.*)'
        target_label: 'container'
      - source_labels: ['__meta_docker_container_log_stream']
        target_label: 'logstream'
      - source_labels: ['__meta_docker_container_label_logging']
        target_label: 'logging'
    pipeline_stages:
      - json:
          expressions:
            timestamp: timestamp
            level: level
            message: message
            service: service
            tenant_id: tenant_id
      - timestamp:
          source: timestamp
          format: RFC3339Nano
      - labels:
          level:
          service:
          tenant_id:
```

EOF

```
log "✅ Monitoring configuration setup completed"
}
```

```
# Database initialization
init_database() {
```

```
log "    [Initialize] Initializing database..."
```

```
# Create database initialization script
cat > ./init-scripts/01-create-extensions.sql << EOF
-- Create necessary extensions
CREATE EXTENSION IF NOT EXISTS "uuid-ossp";
CREATE EXTENSION IF NOT EXISTS "pgcrypto";
CREATE EXTENSION IF NOT EXISTS "pg_stat_statements";
```

```
-- Create core schema
CREATE SCHEMA IF NOT EXISTS core;
```

```
-- Set up connection limits
ALTER SYSTEM SET max_connections = 200;
ALTER SYSTEM SET shared_buffers = '512MB';
ALTER SYSTEM SET effective_cache_size = '2GB';
ALTER SYSTEM SET maintenance_work_mem = '128MB';
ALTER SYSTEM SET checkpoint_completion_target = 0.9;
ALTER SYSTEM SET wal_buffers = '16MB';
ALTER SYSTEM SET default_statistics_target = 100;
ALTER SYSTEM SET random_page_cost = 1.1;
ALTER SYSTEM SET effective_io_concurrency = 200;
```

```
SELECT pg_reload_conf();
EOF
```

```
cat > ./init-scripts/02-create-core-tables.sql << EOF
-- Core tenant management tables
CREATE TABLE IF NOT EXISTS core.tenants (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    subdomain VARCHAR(63) UNIQUE NOT NULL,
    company_name VARCHAR(255) NOT NULL,
    subscription_tier VARCHAR(50) DEFAULT 'starter',
    database_schema VARCHAR(63) NOT NULL,
    status VARCHAR(20) DEFAULT 'active',
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
```

```
-- Etsy Integration
    etsy_shop_id VARCHAR(100),
    etsy_access_token TEXT,
    etsy_refresh_token TEXT,
    etsy_token_expires_at TIMESTAMP,
```

```
-- Billing
stripe_customer_id VARCHAR(100),
billing_email VARCHAR(255),

-- Settings
settings JSONB DEFAULT '{}',

CONSTRAINT valid_subdomain CHECK (subdomain ~ '^[a-z0-9][a-z0-9-]*[a-z0-9]network'
restart: unless-stopped

# API Gateway
api-gateway:
build:
  context: ./services/gateway
  dockerfile: Dockerfile.prod
  container_name: etsy-gateway
environment:
  - REDIS_URL=redis://redis:6379/0
  - DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
  - JWT_SECRET=${JWT_SECRET}
deploy:
resources:
limits:
  memory: 512M
  cpus: '1.0'
reservations:
  memory: 256M
  cpus: '0.5'
labels:
  - "traefik.enable=true"
  - "traefik.http.routers.api-gateway.rule=PathPrefix('/api')"
  - "traefik.http.routers.api-gateway.tls=true"
  - "traefik.http.routers.api-gateway.tls.certresolver=letsencrypt"
depends_on:
  - postgres
  - redis
networks:
  - etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:8080/health"]
  interval: 30s
  timeout: 10s
  retries: 3
```

```
# Authentication Service
auth-service:
  build:
    context: ./services/auth
    dockerfile: Dockerfile.prod
  container_name: etsy-auth
  environment:
    - DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
    - REDIS_URL=redis://redis:6379/1
    - JWT_SECRET=${JWT_SECRET}
    - BCRYPT_ROUNDS=12
  deploy:
    resources:
      limits:
        memory: 1G
        cpus: '1.0'
      reservations:
        memory: 512M
        cpus: '0.5'
    depends_on:
      - postgres
      - redis
  networks:
    - etsy-network
  restart: unless-stopped
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost:8001/health"]
    interval: 30s
    timeout: 10s
    retries: 3
```

```
# Etsy Integration Service
etsy-service:
  build:
    context: ./services/etsy
    dockerfile: Dockerfile.prod
  container_name: etsy-integration
  environment:
    - DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
    - REDIS_URL=redis://redis:6379/2
    - ETSY_API_BASE_URL=https://api.etsy.com/v3
    - ETSY_RATE_LIMIT_CALLS=10000
    - ETSY_RATE_LIMIT_WINDOW=3600
```

```
deploy:
resources:
limits:
  memory: 1G
  cpus: '1.0'
reservations:
  memory: 512M
  cpus: '0.5'
depends_on:
- postgres
- redis
networks:
- etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:8002/health"]
  interval: 30s
  timeout: 10s
  retries: 3
```

```
# Design Management Service
design-service:
build:
context: ./services/design
dockerfile: Dockerfile.prod
container_name: etsy-design
environment:
- DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
- MINIO_ENDPOINT=minio:9000
- MINIO_ACCESS_KEY=${MINIO_ACCESS_KEY}
- MINIO_SECRET_KEY=${MINIO_SECRET_KEY}
- MINIO_SECURE=false
deploy:
resources:
limits:
  memory: 2G
  cpus: '2.0'
reservations:
  memory: 1G
  cpus: '1.0'
depends_on:
- postgres
- minio
networks:
```

```
- etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:8003/health"]
  interval: 30s
  timeout: 10s
  retries: 3

# Analytics Service
analytics-service:
  build:
    context: ./services/analytics
    dockerfile: Dockerfile.prod
    container_name: etsy-analytics
  environment:
    - DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
    - REDIS_URL=redis://redis:6379/3
  deploy:
    resources:
      limits:
        memory: 1G
        cpus: '1.0'
      reservations:
        memory: 512M
        cpus: '0.5'
    depends_on:
      - postgres
      - redis
  networks:
    - etsy-network
restart: unless-stopped
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:8004/health"]
  interval: 30s
  timeout: 10s
  retries: 3

# Job Processing Workers
job-worker:
  build:
    context: ./services/jobs
    dockerfile: Dockerfile.prod
    container_name: etsy-worker
  environment:
```

```
- DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas
- REDIS_URL=redis://redis:6379/4
- WORKER_CONCURRENCY=4
- WORKER_MAX_MEMORY=1G

deploy:
  replicas: 2
  resources:
    limits:
      memory: 1G
      cpus: '1.0'
    reservations:
      memory: 512M
      cpus: '0.5'
  depends_on:
    - redis
    - postgres
  networks:
    - etsy-network
  restart: unless-stopped

# Notification Service
notification-service:
  build:
    context: ./services/notifications
    dockerfile: Dockerfile.prod
  container_name: etsy-notifications
  environment:
    - REDIS_URL=redis://redis:6379/5
  deploy:
    resources:
      limits:
        memory: 512M
        cpus: '0.5'
      reservations:
        memory: 256M
        cpus: '0.25'
    depends_on:
      - redis
  networks:
    - etsy-network
  restart: unless-stopped
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost:8005/health"]
    interval: 30s
```

```

timeout: 10s
retries: 3

# Infrastructure Services
postgres:
  image: postgres:15
  container_name: etsy-postgres
  environment:
    - POSTGRES_DB=etsy_saas
    - POSTGRES_USER=postgres
    - POSTGRES_PASSWORD=${POSTGRES_PASSWORD}
    - POSTGRES_SHARED_BUFFERS=512MB
    - POSTGRES_EFFECTIVE_CACHE_SIZE=2GB
    - POSTGRES_MAINTENANCE_WORK_MEM=128MB
    - POSTGRES_WAL_BUFFERS=16MB
  volumes:
    - postgres_data:/var/lib/postgresql/data
    - ./init-scripts:/docker-entrypoint-initdb.d
    - ./postgres-config:/etc/postgresql/postgresql.conf
  ports:
    - "5432:5432"
deploy:
  resources:
    limits:
      memory: 3G
      cpus: '2.0'
    reservations:
      memory: 2G
      cpus: '1.0'
networks:
  - etsy-# Multi-Tenant Etsy Seller Automater Backend Architecture Deep Dive

```

Executive Summary

This document outlines the backend architecture **for** transforming the existing Etsy Seller Automater into a scalab

Current State Analysis

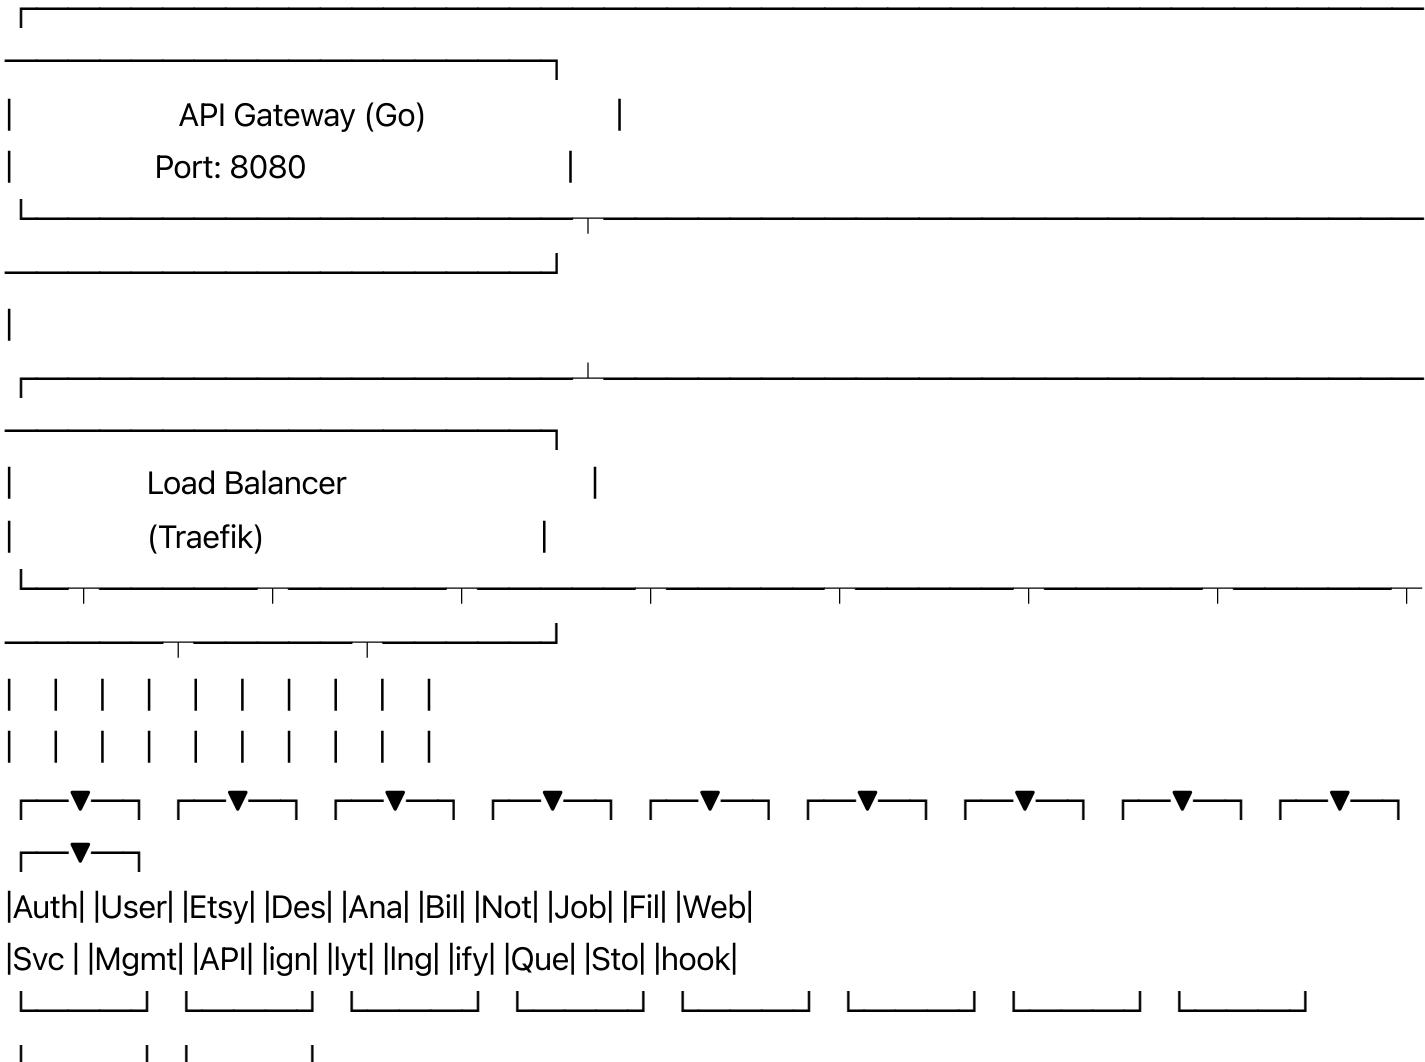
Existing Architecture

- **Framework**: FastAPI (Python)
- **Database**: PostgreSQL
- **Authentication**: OAuth 2.0 with PKCE **for** Etsy API
- **Frontend**: React with Tailwind CSS
- **Deployment**: Docker Compose

- **Key Features**:
 - Etsy shop analytics
 - Design management and mockup creation
 - Listing automation
 - Mask creator for product images

Proposed Multi-Tenant Architecture

1. Service Architecture Overview



2. Core Services Architecture

2.1 API Gateway Service (Go)

Port: 8080

Responsibilities: Request routing, rate limiting, authentication middleware, tenant isolation

```
```go
// pkg/gateway/main.go
package main

import (
 "context"
 "log"
 "net/http"
 "time"

 "github.com/gin-gonic/gin"
 "github.com/redis/go-redis/v9"
)

type Gateway struct {
 router *gin.Engine
 redisClient *redis.Client
 services map[string]string
}

type TenantMiddleware struct {
 redis *redis.Client
}

func (tm *TenantMiddleware) ExtractTenant() gin.HandlerFunc {
 return func(c *gin.Context) {
 host := c.Request.Host
 subdomain := extractSubdomain(host)

 // Validate tenant exists
 tenantID, err := tm.redis.Get(c.Request.Context(),
 fmt.Sprintf("tenant:%s", subdomain)).Result()
 if err != nil {
 c.JSON(404, gin.H{"error": "Tenant not found"})
 c.Abort()
 return
 }
 }
}
```

```
}

c.Set("tenant_id", tenantID)
c.Set("subdomain", subdomain)
c.Next()
}

}

func (g *Gateway) setupRoutes() {
 v1 := g.router.Group("/api/v1")
 v1.Use(g.tenantMiddleware.ExtractTenant())

 // Route to microservices
 v1.Any("/auth/*path", g.proxyToService("auth-service"))
 v1.Any("/etsy/*path", g.proxyToService("etsy-service"))
 v1.Any("/designs/*path", g.proxyToService("design-service"))
 v1.Any("/analytics/*path", g.proxyToService("analytics-service"))
}
```

## 2.2 Authentication & User Management Service (Python/FastAPI)

**Port:** 8001

**Database:** `tenant_users` schema in PostgreSQL

python

```
services/auth/main.py
from fastapi import FastAPI, Depends, HTTPException
from sqlalchemy.orm import Session
import jwt
from datetime import datetime, timedelta

app = FastAPI(title="Authentication Service")

class AuthService:
 def __init__(self, db: Session):
 self.db = db

 async def authenticate_tenant_user(self,
 tenant_id: str,
 email: str,
 password: str) -> dict:
 """Authenticate user within specific tenant context"""
 user = self.db.query(TenantUser).filter(
 TenantUser.tenant_id == tenant_id,
 TenantUser.email == email,
 TenantUser.is_active == True
).first()

 if not user or not verify_password(password, user.hashed_password):
 raise HTTPException(401, "Invalid credentials")

 # Generate tenant-scoped JWT
 payload = {
 "user_id": user.id,
 "tenant_id": tenant_id,
 "permissions": user.permissions,
 "exp": datetime.utcnow() + timedelta(hours=24)
 }

 token = jwt.encode(payload, JWT_SECRET, algorithm="HS256")

 return {
 "access_token": token,
 "token_type": "bearer",
 "user": user.to_dict(),
 "tenant": user.tenant.to_dict()
 }
```

```

Database Models
class Tenant(Base):
 __tablename__ = "tenants"

 id = Column(UUID(as_uuid=True), primary_key=True, default=uuid.uuid4)
 subdomain = Column(String(63), unique=True, nullable=False)
 company_name = Column(String(255), nullable=False)
 subscription_tier = Column(String(50), default="basic")
 created_at = Column(DateTime, default=datetime.utcnow)

 # Etsy Integration
 etsy_shop_id = Column(String(100))
 etsy_access_token = Column(Text)
 etsy_refresh_token = Column(Text)

class TenantUser(Base):
 __tablename__ = "tenant_users"

 id = Column(UUID(as_uuid=True), primary_key=True, default=uuid.uuid4)
 tenant_id = Column(UUID(as_uuid=True), ForeignKey("tenants.id"))
 email = Column(String(255), nullable=False)
 hashed_password = Column(String(255), nullable=False)
 permissions = Column(JSON) # Role-based permissions
 is_active = Column(Boolean, default=True)

```

## 2.3 Etsy Integration Service (Python/FastAPI)

**Port:** 8002

**Focus:** OAuth flows, API calls, rate limiting per tenant

python

```
services/etsy/main.py
import asyncio
import aiohttp
from fastapi import FastAPI, BackgroundTasks
from rq import Queue
import redis

app = FastAPI(title="Etsy Integration Service")

class EtsyService:
 def __init__(self, tenant_id: str):
 self.tenant_id = tenant_id
 self.base_url = "https://api.etsy.com/v3"

 @async def get_shop_analytics(self, date_range: str) -> dict:
 """Get shop analytics for tenant with rate limiting"""

 # Implement tenant-specific rate limiting
 rate_limit_key = f"etsy_rate_limit:{self.tenant_id}"

 async with aiohttp.ClientSession() as session:
 headers = await self._get_auth_headers()

 # Get receipt data
 receipts = await self._fetch_receipts(session, headers, date_range)

 # Process analytics
 analytics = await self._process_analytics_data(receipts)

 return {
 "tenant_id": self.tenant_id,
 "date_range": date_range,
 "total_revenue": analytics["revenue"],
 "total_orders": analytics["orders"],
 "top_products": analytics["top_products"],
 "monthly_breakdown": analytics["monthly"]
 }

 @async def create_listing(self, listing_data: dict) -> dict:
 """Create Etsy listing with tenant isolation"""

 # Validate listing data against tenant's subscription limits
 await self._validate_listing_limits()
```

```

Create listing via Etsy API
async with aiohttp.ClientSession() as session:
 headers = await self._get_auth_headers()

 response = await session.post(
 f"{self.base_url}/application/shops/{self.shop_id}/listings",
 headers=headers,
 json=listing_data
)

 result = await response.json()

Store listing in tenant database
await self._store_listing_record(result)

return result

Background job for bulk operations
@app.post("/bulk-listing-creation")
async def bulk_create_listings(
 listings: List[dict],
 tenant_id: str,
 background_tasks: BackgroundTasks
):
 """Queue bulk listing creation as background job"""

 redis_conn = redis.Redis(host='redis', port=6379, db=0)
 queue = Queue('bulk_operations', connection=redis_conn)

 job = queue.enqueue(
 'tasks.bulk_create_listings',
 listings,
 tenant_id,
 job_timeout='30m'
)

 return {"job_id": job.id, "status": "queued"}

```

## 2.4 Design Management Service (Go)

**Port:** 8003

**Focus:** File storage, image processing, mockup generation

go

```
// services/design/main.go
package main

import (
 "context"
 "fmt"
 "io"
 "path/filepath"

 "github.com/gin-gonic/gin"
 "github.com/minio/minio-go/v7"
 "gorm.io/gorm"
)

type DesignService struct {
 db *gorm.DB
 minioClient *minio.Client
 router *gin.Engine
}

type Design struct {
 ID uint `gorm:"primaryKey"`
 TenantID string `gorm:"index"`
 Name string
 FilePath string
 FileSize int64
 ContentType string
 Tags []string `gorm:"serializer:json"`
 CreatedAt time.Time
}

func (ds *DesignService) uploadDesign(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 file, header, err := c.Request.FormFile("design")
 if err != nil {
 c.JSON(400, gin.H{"error": "Invalid file"})
 return
 }
 defer file.Close()

 // Generate tenant-scoped file path
 fileName := fmt.Sprintf("%s/%s/%s",
 tenantID,
 header.Filename,
 file.Size)
}
```

```
tenantID, "designs", header.Filename)

// Upload to MinIO
_, err = ds.minioClient.PutObject(
 context.Background(),
 "tenant-designs",
 fileName,
 file,
 header.Size,
 minio.PutObjectOptions{
 ContentType: header.Header.Get("Content-Type"),
 },
)

if err != nil {
 c.JSON(500, gin.H{"error": "Upload failed"})
 return
}

// Save metadata to database
design := Design{
 TenantID: tenantID,
 Name: header.Filename,
 FilePath: fileName,
 FileSize: header.Size,
 ContentType: header.Header.Get("Content-Type"),
}
ds.db.Create(&design)

c.JSON(201, design)
}

func (ds *DesignService) generateMockup(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
 designID := c.Param("id")

 i// Get design from database
 var design Design
 result := ds.db.Where("tenant_id = ? AND id = ?",
 tenantID, designID).First(&design)

 if result.Error != nil {
 c.JSON(404, gin.H{"error": "Design not found"})
 }
}
```

```
 return
}

// Queue mockup generation
job := MockupGenerationJob{
 TenantID: tenantID,
 DesignID: designID,
 FilePath: design.FilePath,
}

// Submit to job queue (Redis/RQ)
jobID, err := ds.queueMockupJob(job)
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to queue job"})
 return
}

c.JSON(202, gin.H{
 "job_id": jobID,
 "status": "queued",
 "message": "Mockup generation started"
})
}
```

## 2.5 Analytics Service (Python/FastAPI)

**Port:** 8004

**Focus:** Data aggregation, reporting, business intelligence

python



```

analytics = {
 "summary": {
 "total_revenue": df['daily_revenue'].sum(),
 "total_orders": df['daily_orders'].sum(),
 "avg_daily_revenue": df['daily_revenue'].mean(),
 "revenue_growth": self._calculate_growth_rate(df)
 },
 "daily_breakdown": df.to_dict('records'),
 "top_performing_days": df.nlargest(5, 'daily_revenue').to_dict('records')
}

return analytics

```

```

async def get_product_performance(self) -> dict:
 """Analyze product performance metrics"""

 query = text("""
SELECT
 p.name,
 p.etsy_listing_id,
 COUNT(oi.id) as total_sales,
 SUM(oi.quantity) as units_sold,
 SUM(oi.price * oi.quantity) as total_revenue,
 AVG(oi.price) as avg_price
FROM products p
LEFT JOIN order_items oi ON p.id = oi.product_id
LEFT JOIN tenant_orders o ON oi.order_id = o.id
WHERE o.tenant_id = :tenant_id
GROUP BY p.id, p.name, p.etsy_listing_id
ORDER BY total_revenue DESC
 """)

```

```

result = self.db.execute(query, {"tenant_id": self.tenant_id})

```

```

products = []
for row in result:
 products.append({
 "name": row.name,
 "etsy_listing_id": row.etsy_listing_id,
 "total_sales": row.total_sales,
 "units_sold": row.units_sold,
 "total_revenue": float(row.total_revenue or 0),
 "avg_price": float(row.avg_price or 0)
 })

```

```

return {
 "top_products": products[:10],
 "total_products": len(products),
 "products_with_sales": len([p for p in products if p['total_sales'] > 0])
}

@app.get("/analytics/dashboard")
async def get_dashboard_data(
 tenant_id: str = Depends(get_tenant_id),
 date_range: str = "30d"
):
 """Get comprehensive dashboard analytics"""

 analytics_service = AnalyticsService(tenant_id, get_db())

 # Calculate date range
 end_date = datetime.utcnow()
 days = int(date_range.replace('d', ''))
 start_date = end_date - timedelta(days=days)

 sales_data = await analytics_service.get_sales_analytics(start_date, end_date)
 product_data = await analytics_service.get_product_performance()

 return {
 "tenant_id": tenant_id,
 "date_range": date_range,
 "sales": sales_data,
 "products": product_data,
 "generated_at": datetime.utcnow().isoformat()
}

```

### 3. Database Architecture

#### 3.1 Multi-Tenant Schema Design (PostgreSQL)

sql

```

-- Core tenant management
CREATE SCHEMA IF NOT EXISTS core;

CREATE TABLE core.tenants (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 subdomain VARCHAR(63) UNIQUE NOT NULL,
 company_name VARCHAR(255) NOT NULL,
 subscription_tier VARCHAR(50) DEFAULT 'basic',
 database_schema VARCHAR(63) NOT NULL,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE core.tenant_users (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id),
 email VARCHAR(255) NOT NULL,
 hashed_password VARCHAR(255) NOT NULL,
 role VARCHAR(50) DEFAULT 'user',
 permissions JSONB,
 is_active BOOLEAN DEFAULT TRUE,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Function to create tenant schema
CREATE OR REPLACE FUNCTION create_tenant_schema(tenant_id UUID, schema_name VARCHAR)
RETURNS VOID AS $$
BEGIN
 -- Create schema
 EXECUTE format('CREATE SCHEMA IF NOT EXISTS %I', schema_name);

 -- Create tenant-specific tables
 EXECUTE format(
 CREATE TABLE %I.orders (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 etsy_receipt_id BIGINT,
 total_amount DECIMAL(10,2),
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
), schema_name);

 EXECUTE format(
 CREATE TABLE %I.products (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```
 name VARCHAR(255) NOT NULL,
 etsy_listing_id BIGINT,
 price DECIMAL(10,2),
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)', schema_name);
```

```
EXECUTE format('
```

```
CREATE TABLE %I.designs (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 name VARCHAR(255) NOT NULL,
 file_path VARCHAR(500),
 file_size BIGINT,
 tags TEXT[],
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)', schema_name);
```

```
END;
```

```
$$ LANGUAGE plpgsql;
```

## 4. Container Orchestration (Docker Compose)

```
yaml
```

```
docker-compose.yml
version: '3.8'

services:
API Gateway
api-gateway:
build:
 context: ./services/gateway
 dockerfile: Dockerfile
ports:
- "8080:8080"
environment:
- REDIS_URL=redis://redis:6379
depends_on:
- redis
networks:
- etsy-network

Authentication Service
auth-service:
build:
 context: ./services/auth
 dockerfile: Dockerfile
ports:
- "8001:8001"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- JWT_SECRET=your-jwt-secret-here
- REDIS_URL=redis://redis:6379
depends_on:
- postgres
- redis
networks:
- etsy-network

Etsy Integration Service
etsy-service:
build:
 context: ./services/etsy
 dockerfile: Dockerfile
ports:
- "8002:8002"
environment:
```

```
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- REDIS_URL=redis://redis:6379
- ETSY_API_BASE_URL=https://api.etsy.com/v3

depends_on:
- postgres
- redis

networks:
- etsy-network

Design Management Service (Go)
design-service:
build:
context: ./services/design
dockerfile: Dockerfile
ports:
- "8003:8003"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- MINIO_ENDPOINT=minio:9000
- MINIO_ACCESS_KEY=minioadmin
- MINIO_SECRET_KEY=minioadmin
depends_on:
- postgres
- minio
networks:
- etsy-network

Analytics Service
analytics-service:
build:
context: ./services/analytics
dockerfile: Dockerfile
ports:
- "8004:8004"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
depends_on:
- postgres
networks:
- etsy-network

Job Queue Worker
job-worker:
build:
```

```
context: ./services/jobs
dockerfile: Dockerfile
environment:
 - REDIS_URL=redis://redis:6379
 - DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
depends_on:
 - redis
 - postgres
networks:
 - etsy-network

Infrastructure Services

postgres:
 image: postgres:15
 environment:
 - POSTGRES_DB=etsy_saas
 - POSTGRES_USER=postgres
 - POSTGRES_PASSWORD=password
 volumes:
 - postgres_data:/var/lib/postgresql/data
 - ./init.sql:/docker-entrypoint-initdb.d/init.sql
 ports:
 - "5432:5432"
 networks:
 - etsy-network

redis:
 image: redis:7-alpine
 ports:
 - "6379:6379"
 volumes:
 - redis_data:/data
 networks:
 - etsy-network

minio:
 image: minio/minio
 ports:
 - "9000:9000"
 - "9001:9001"
 volumes:
 - minio_data:/data
 environment:
 - MINIO_ROOT_USER=minioadmin
```

```
- MINIO_ROOT_PASSWORD=minioadmin
command: server /data --console-address ":9001"
networks:
- etsy-network

Reverse Proxy
traefik:
image: traefik:v2.10
command:
- --providers.docker=true
- --providers.docker.exposedbydefault=false
- --entrypoints.web.address=:80
- --entrypoints.websecure.address=:443
ports:
- "80:80"
- "443:443"
- "8080:8080" # Dashboard
volumes:
- /var/run/docker.sock:/var/run/docker.sock:ro
networks:
- etsy-network

volumes:
postgres_data:
redis_data:
minio_data:

networks:
etsy-network:
driver: bridge
```

## 5. Deployment on QNAP NAS

### 5.1 QNAP-Specific Configuration

```
bash
```

```
#!/bin/bash
deploy-qnap.sh

Install Container Station if not already installed
Access QNAP Web UI -> App Center -> Container Station

Create project directory
ssh admin@qnap-ip "mkdir -p /share/Container/etsy-saas"

Copy docker-compose files
scp docker-compose.yml admin@qnap-ip:/share/Container/etsy-saas/
scp docker-compose.override.yml admin@qnap-ip:/share/Container/etsy-saas/

Create persistent volumes on QNAP
ssh admin@qnap-ip "mkdir -p /share/Container/volumes/{postgres,redis,minio}"

Set proper permissions
ssh admin@qnap-ip "chmod -R 755 /share/Container/etsy-saas"

Start services
ssh admin@qnap-ip "cd /share/Container/etsy-saas && docker-compose up -d"
```

## 5.2 QNAP Resource Optimization

yaml

```
docker-compose.override.yml - QNAP specific optimizations
version: '3.8'

services:
 api-gateway:
 deploy:
 resources:
 limits:
 memory: 256M
 cpus: '0.5'

 auth-service:
 deploy:
 resources:
 limits:
 memory: 512M
 cpus: '0.5'

 etsy-service:
 deploy:
 resources:
 limits:
 memory: 512M
 cpus: '0.5'

 design-service:
 deploy:
 resources:
 limits:
 memory: 1G
 cpus: '1.0'

 postgres:
 deploy:
 resources:
 limits:
 memory: 2G
 cpus: '1.0'

 environment:
 - POSTGRES_SHARED_BUFFERS=256MB
 - POSTGRES_EFFECTIVE_CACHE_SIZE=1GB
```

## 6. Monitoring and Observability

### 6.1 Health Checks

```
python
```

```
services/common/health.py
from fastapi import FastAPI
import asyncio
import psycopg2
import redis

class HealthCheck:
 def __init__(self, app: FastAPI):
 self.app = app
 self.setup_routes()

 def setup_routes(self):
 @self.app.get("/health")
 async def health_check():
 checks = {
 "database": await self.check_database(),
 "redis": await self.check_redis(),
 "service": "healthy"
 }

 all_healthy = all(
 status == "healthy"
 for status in checks.values()
)

 return {
 "status": "healthy" if all_healthy else "unhealthy",
 "checks": checks
 }

 async def check_database(self) -> str:
 try:
 conn = psycopg2.connect(DATABASE_URL)
 conn.close()
 return "healthy"
 except Exception as e:
 return f"unhealthy: {str(e)}"

 async def check_redis(self) -> str:
 try:
 r = redis.Redis.from_url(REDIS_URL)
 r.ping()
 return "healthy"
```

```
except Exception as e:
 return f"unhealthy: {str(e)}"
```

## 7. Security Considerations

### 7.1 Tenant Isolation Security

python

```
services/common/security.py
from functools import wraps
import jwt
from fastapi import HTTPException

def require_tenant_access(required_permissions: List[str] = None):
 def decorator(func):
 @wraps(func)
 async def wrapper(*args, **kwargs):
 # Extract JWT token
 token = extract_token_from_request()

 try:
 payload = jwt.decode(token, JWT_SECRET, algorithms=["HS256"])
 tenant_id = payload.get("tenant_id")
 user_permissions = payload.get("permissions", [])

 # Verify tenant access
 if not tenant_id:
 raise HTTPException(403, "No tenant access")

 # Check specific permissions if required
 if required_permissions:
 if not any(perm in user_permissions for perm in required_permissions):
 raise HTTPException(403, "Insufficient permissions")

 # Add to request context
 kwargs['tenant_id'] = tenant_id
 kwargs['user_permissions'] = user_permissions

 except jwt.InvalidTokenError:
 raise HTTPException(401, "Invalid token")

 return await func(*args, **kwargs)
 return wrapper
 return decorator

Usage example
@app.post("/designs/upload")
@require_tenant_access(["design:create"])
async def upload_design(tenant_id: str, user_permissions: List[str], file: UploadFile = File(...)):
 """Upload design with tenant isolation"""

```

```
Process upload with tenant context
```

```
pass
```

## 8. Background Job Processing Architecture

### 8.1 Job Queue System (Python/RQ + Redis)

```
python
```

```
services/jobs/worker.py
from rq import Worker, Queue, Connection
import redis
from typing import Dict, Any
import logging

Job definitions
def process_bulk_listings(tenant_id: str, listings_data: List[Dict]) -> Dict[str, Any]:
 """Process bulk listing creation for a tenant"""
 results = {
 "tenant_id": tenant_id,
 "total_listings": len(listings_data),
 "successful": 0,
 "failed": 0,
 "errors": []
 }

 etsy_service = EtsyService(tenant_id)

 for listing_data in listings_data:
 try:
 result = await etsy_service.create_listing(listing_data)
 results["successful"] += 1

 # Store success in tenant database
 await store_listing_result(tenant_id, listing_data["id"], result)

 except Exception as e:
 results["failed"] += 1
 results["errors"].append({
 "listing_id": listing_data.get("id"),
 "error": str(e)
 })

 # Log error for tenant
 await log_tenant_error(tenant_id, "listing_creation", str(e))

 return results

def generate_mockup_images(tenant_id: str, design_id: str, mockup_config: Dict) -> Dict[str, Any]:
 """Generate mockup images for a design"""
 import PIL.Image
 import numpy as np
```

```
try:
 # Load design from MinIO
 design_service = DesignService(tenant_id)
 design_data = design_service.get_design(design_id)

 # Load base mockup templates
 mockup_templates = load_mockup_templates(mockup_config["template_type"])

 generated_mockups = []

 for template in mockup_templates:
 # Apply design to template
 mockup_image = apply_design_to_template(
 design_data["image_data"],
 template,
 mockup_config.get("mask_data", {}))
)

 # Save mockup to MinIO
 mockup_path = f"{tenant_id}/mockups/{design_id}_{template['name']}.png"
 mockup_url = design_service.save_mockup(mockup_path, mockup_image)

 generated_mockups.append({
 "template_name": template["name"],
 "mockup_url": mockup_url,
 "size": mockup_image.size
 })

 return {
 "tenant_id": tenant_id,
 "design_id": design_id,
 "generated_mockups": generated_mockups,
 "status": "completed"
 }

except Exception as e:
 logging.error(f"Mockup generation failed for {tenant_id}:{design_id}: {e}")
 return {
 "tenant_id": tenant_id,
 "design_id": design_id,
 "status": "failed",
 "error": str(e)
 }
```

```
def sync_etsy_data(tenant_id: str) -> Dict[str, Any]:
 """Sync Etsy data for tenant (orders, listings, etc.)"""
 etsy_service = EtsyService(tenant_id)

 # Get latest sync timestamp
 last_sync = get_last_sync_timestamp(tenant_id)

 sync_results = {
 "tenant_id": tenant_id,
 "sync_started": datetime.utcnow().isoformat(),
 "orders_synced": 0,
 "listings_synced": 0,
 "errors": []
 }

 try:
 # Sync orders
 new_orders = await etsy_service.get_orders_since(last_sync)
 for order in new_orders:
 await store_tenant_order(tenant_id, order)
 sync_results["orders_synced"] += 1

 # Sync listings
 updated_listings = await etsy_service.get_listings_updates_since(last_sync)
 for listing in updated_listings:
 await store_tenant_listing(tenant_id, listing)
 sync_results["listings_synced"] += 1

 # Update sync timestamp
 await update_sync_timestamp(tenant_id, datetime.utcnow())

 sync_results["status"] = "completed"
 sync_results["sync_completed"] = datetime.utcnow().isoformat()

 except Exception as e:
 sync_results["status"] = "failed"
 sync_results["errors"].append(str(e))
 logging.error(f"Etsy sync failed for {tenant_id}: {e}")

 return sync_results

Worker setup
if __name__ == '__main__':
```

```
redis_conn = redis.from_url(REDIS_URL)

Create queues with priorities
high_priority_queue = Queue('high_priority', connection=redis_conn)
normal_queue = Queue('normal', connection=redis_conn)
low_priority_queue = Queue('low_priority', connection=redis_conn)

Start worker
worker = Worker([high_priority_queue, normal_queue, low_priority_queue],
 connection=redis_conn)
worker.work()
```

## 8.2 Job Management API (Go)

```
go
```

```
// services/jobs/main.go
package main

import (
 "encoding/json"
 "net/http"
 "time"

 "github.com/gin-gonic/gin"
 "github.com/go-redis/redis/v8"
)

type JobService struct {
 redis *redis.Client
 router *gin.Engine
}

type Job struct {
 ID string `json:"id"`
 TenantID string `json:"tenant_id"`
 Type string `json:"type"`
 Status string `json:"status"`
 Payload map[string]interface{} `json:"payload"`
 Result map[string]interface{} `json:"result,omitempty"`
 CreatedAt time.Time `json:"created_at"`
 CompletedAt *time.Time `json:"completed_at,omitempty"`
 Error string `json:"error,omitempty"`
}

func (js *JobService) queueJob(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 var jobRequest struct {
 Type string `json:"type" binding:"required"`
 Payload map[string]interface{} `json:"payload" binding:"required"`
 Priority string `json:"priority" // high, normal, low`
 }

 if err := c.ShouldBindJSON(&jobRequest); err != nil {
 c.JSON(400, gin.H{"error": err.Error()})
 return
 }
```

```

// Generate job ID
jobID := generateJobID()

job := Job{
 ID: jobID,
 TenantID: tenantID,
 Type: jobRequest.Type,
 Status: "queued",
 Payload: jobRequest.Payload,
 CreatedAt: time.Now(),
}

// Store job metadata
jobData, _ := json.Marshal(job)
js.redis.Set(c.Request.Context(),
 fmt.Sprintf("job:%s", jobID), jobData, 24*time.Hour)

// Queue job based on type and priority
queueName := js.getQueueForJob(jobRequest.Type, jobRequest.Priority)

jobPayload := map[string]interface{}{
 "job_id": jobID,
 "tenant_id": tenantID,
 "type": jobRequest.Type,
 "payload": jobRequest.Payload,
}

payloadJSON, _ := json.Marshal(jobPayload)

err := js.redis.LPush(c.Request.Context(), queueName, payloadJSON).Err()
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to queue job"})
 return
}

c.JSON(202, gin.H{
 "job_id": jobID,
 "status": "queued",
 "message": fmt.Sprintf("Job queued in %s", queueName)
})

func (js *JobService) getJobStatus(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
}

```

```
jobID := c.Param("job_id")

// Get job data
jobData, err := js.redis.Get(c.Request.Context()),
 fmt.Sprintf("job:%s", jobID)).Result()
if err == redis.Nil {
 c.JSON(404, gin.H{"error": "Job not found"})
 return
}

var job Job
json.Unmarshal([]byte(jobData), &job)

// Verify tenant access
if job.TenantID != tenantID {
 c.JSON(403, gin.H{"error": "Access denied"})
 return
}

c.JSON(200, job)
}

func (js *JobService) listTenantJobs(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 // Get all job keys for tenant
pattern := fmt.Sprintf("job:*")
keys, err := js.redis.Keys(c.Request.Context(), pattern).Result()
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to fetch jobs"})
 return
}

var jobs []Job
for _, key := range keys {
 jobData, err := js.redis.Get(c.Request.Context(), key).Result()
 if err != nil {
 continue
 }

 var job Job
 if err := json.Unmarshal([]byte(jobData), &job); err != nil {
 continue
 }
}
```

```
// Filter by tenant
if job.TenantID == tenantID {
 jobs = append(jobs, job)
}
}

c.JSON(200, gin.H{
 "jobs": jobs,
 "total": len(jobs)
})
}
```

## 9. Advanced Features Implementation

### 9.1 Real-time Notifications Service (Go + WebSockets)

```
go
```

```
// services/notifications/main.go
package main

import (
 "encoding/json"
 "log"
 "net/http"

 "github.com/gin-gonic/gin"
 "github.com/gorilla/websocket"
)

type NotificationService struct {
 clients map[string]map[*websocket.Conn]bool // tenant_id -> connections
 broadcast chan NotificationMessage
 register chan *Client
 unregister chan *Client
 router *gin.Engine
}

type Client struct {
 TenantID string
 Conn *websocket.Conn
 Send chan NotificationMessage
}

type NotificationMessage struct {
 TenantID string `json:"tenant_id"`
 Type string `json:"type"`
 Title string `json:"title"`
 Message string `json:"message"`
 Data interface{} `json:"data,omitempty"`
 Timestamp int64 `json:"timestamp"`
}

var upgrader = websocket.Upgrader{
 CheckOrigin: func(r *http.Request) bool {
 return true // Configure proper origin checking in production
 },
}

func (ns *NotificationService) handleWebSocket(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
```

```
conn, err := upgrader.Upgrade(c.Writer, c.Request, nil)
if err != nil {
 log.Printf("WebSocket upgrade failed: %v", err)
 return
}

client := &Client{
 TenantID: tenantID,
 Conn: conn,
 Send: make(chan NotificationMessage, 256),
}
ns.register <- client

go client.writePump()
go client.readPump(ns.unregister)
}

func (c *Client) writePump() {
 defer c.Conn.Close()

 for {
 select {
 case message, ok := <-c.Send:
 if !ok {
 c.Conn.WriteMessage(websocket.CloseMessage, []byte{})
 return
 }

 if err := c.Conn.WriteJSON(message); err != nil {
 log.Printf("WebSocket write error: %v", err)
 return
 }
 }
 }
}

func (ns *NotificationService) run() {
 for {
 select {
 case client := <-ns.register:
 if ns.clients[client.TenantID] == nil {
 ns.clients[client.TenantID] = make(map[*websocket.Conn]bool)
 }
 }
 }
}
```

```

 }

 ns.clients[client.TenantID][client.Conn] = true
 log.Printf("Client connected for tenant: %s", client.TenantID)

 case client := <-ns.unregister:
 if clients, ok := ns.clients[client.TenantID]; ok {
 if _, ok := clients[client.Conn]; ok {
 delete(clients, client.Conn)
 close(client.Send)
 client.Conn.Close()
 }
 }
 }

 case message := <-ns.broadcast:
 if clients, ok := ns.clients[message.TenantID]; ok {
 for conn := range clients {
 select {
 case conn.WriteJSON(message):
 default:
 close(conn.Send)
 delete(clients, conn)
 conn.Close()
 }
 }
 }
 }

}

// Notification triggers
func (ns *NotificationService) sendJobCompletion(tenantID, jobID, jobType string, result interface{}) {
 message := NotificationMessage{
 TenantID: tenantID,
 Type: "job_completed",
 Title: "Job Completed",
 Message: fmt.Sprintf("%s job completed successfully", jobType),
 Data: map[string]interface{}{"job_id": jobID, "result": result},
 Timestamp: time.Now().Unix(),
 }

 ns.broadcast <- message
}

```

## 9.2 Billing and Subscription Management (Python)

python

```
services/billing/main.py
from fastapi import FastAPI, Depends, HTTPException
from sqlalchemy.orm import Session
from datetime import datetime, timedelta
from typing import Dict, Any
import stripe

app = FastAPI(title="Billing Service")

Subscription tiers configuration
SUBSCRIPTION_TIERS = {
 "starter": {
 "name": "Starter",
 "monthly_price": 29,
 "limits": {
 "max_listings": 100,
 "max_designs": 50,
 "max_mockups_per_month": 500,
 "api_calls_per_hour": 100
 }
 },
 "professional": {
 "name": "Professional",
 "monthly_price": 79,
 "limits": {
 "max_listings": 1000,
 "max_designs": 500,
 "max_mockups_per_month": 5000,
 "api_calls_per_hour": 500
 }
 },
 "enterprise": {
 "name": "Enterprise",
 "monthly_price": 199,
 "limits": {
 "max_listings": -1, # Unlimited
 "max_designs": -1, # Unlimited
 "max_mockups_per_month": -1, # Unlimited
 "api_calls_per_hour": 2000
 }
 }
}
```

```
class BillingService:
 def __init__(self, tenant_id: str, db: Session):
 self.tenant_id = tenant_id
 self.db = db

 @async def get_usage_stats(self) -> Dict[str, Any]:
 """Get current usage statistics for tenant"""

 # Query current month usage
 current_month = datetime.now().replace(day=1, hour=0, minute=0, second=0)

 usage_stats = {
 "current_listings": await self._count_active_listings(),
 "total_designs": await self._count_designs(),
 "mockups_this_month": await self._count_mockups_since(current_month),
 "api_calls_this_hour": await self._count_api_calls_last_hour()
 }

 # Get subscription limits
 subscription = await self._get_tenant_subscription()
 tier_limits = SUBSCRIPTION_TIERS[subscription.tier]["limits"]

 # Calculate usage percentages
 usage_percentages = {}
 for metric, current_value in usage_stats.items():
 limit_key = f"max_{metric}" if metric != "api_calls_this_hour" else "api_calls_per_hour"
 limit = tier_limits.get(limit_key, -1)

 if limit == -1: # Unlimited
 usage_percentages[metric] = 0
 else:
 usage_percentages[metric] = (current_value / limit) * 100

 return {
 "tenant_id": self.tenant_id,
 "subscription_tier": subscription.tier,
 "usage": usage_stats,
 "limits": tier_limits,
 "usage_percentages": usage_percentages,
 "billing_period": {
 "start": subscription.current_period_start,
 "end": subscription.current_period_end
 }
 }
```

```

async def check_usage_limits(self, resource_type: str, requested_amount: int = 1) -> bool:
 """Check if tenant can use more resources"""

 usage_stats = await self.get_usage_stats()
 current_usage = usage_stats["usage"]
 limits = usage_stats["limits"]

 resource_mapping = {
 "listings": ("current_listings", "max_listings"),
 "designs": ("total_designs", "max_designs"),
 "mockups": ("mockups_this_month", "max_mockups_per_month"),
 "api_calls": ("api_calls_this_hour", "api_calls_per_hour")
 }

 if resource_type not in resource_mapping:
 raise ValueError(f"Unknown resource type: {resource_type}")

 usage_key, limit_key = resource_mapping[resource_type]
 current_value = current_usage[usage_key]
 limit = limits[limit_key]

 if limit == -1: # Unlimited
 return True

 return (current_value + requested_amount) <= limit

async def create_stripe_checkout_session(self,
 new_tier: str,
 success_url: str,
 cancel_url: str) -> str:
 """Create Stripe checkout session for subscription upgrade/downgrade"""

 if new_tier not in SUBSCRIPTION_TIERS:
 raise HTTPException(400, f"Invalid subscription tier: {new_tier}")

 tenant = await self._get_tenant()
 price_amount = SUBSCRIPTION_TIERS[new_tier]["monthly_price"] * 100 # Stripe uses cents

 checkout_session = stripe.checkout.Session.create(
 customer_email=tenant.primary_email,
 payment_method_types=['card'],
 line_items=[{
 'price_data': {

```

```

'currency': 'usd',
'unit_amount': price_amount,
'product_data': {
 'name': f'Etsy Seller Automater - {SUBSCRIPTION_TIERS[new_tier]["name"]}',
 'description': f'Monthly subscription to {new_tier} tier'
},
'recurring': {
 'interval': 'month'
},
'quantity': 1,
}],
mode='subscription',
success_url=success_url + '?session_id={CHECKOUT_SESSION_ID}',
cancel_url=cancel_url,
metadata={
 'tenant_id': self.tenant_id,
 'subscription_tier': new_tier
}
)

return checkout_session.url

@app.get("/billing/usage")
async def get_tenant_usage(tenant_id: str = Depends(get_tenant_id)):
 """Get detailed usage statistics for tenant"""
 billing_service = BillingService(tenant_id, get_db())
 return await billing_service.get_usage_stats()

@app.post("/billing/upgrade")
async def create_upgrade_session(
 new_tier: str,
 tenant_id: str = Depends(get_tenant_id)
):
 """Create Stripe checkout session for subscription upgrade"""
 billing_service = BillingService(tenant_id, get_db())

 checkout_url = await billing_service.create_stripe_checkout_session(
 new_tier=new_tier,
 success_url=f"https://{{tenant_id}}.yourdomain.com/billing/success",
 cancel_url=f"https://{{tenant_id}}.yourdomain.com/billing/cancel"
)

 return {"checkout_url": checkout_url}

```

```

@app.post("/billing/webhook")
async def stripe_webhook(request: Request):
 """Handle Stripe webhooks for subscription events"""
 payload = await request.body()
 sig_header = request.headers.get('Stripe-Signature')

 try:
 event = stripe.Webhook.construct_event(
 payload, sig_header, STRIPE_WEBHOOK_SECRET
)
 except ValueError:
 raise HTTPException(400, "Invalid payload")
 except stripe.error.SignatureVerificationError:
 raise HTTPException(400, "Invalid signature")

 if event['type'] == 'checkout.session.completed':
 session = event['data']['object']
 tenant_id = session['metadata']['tenant_id']
 new_tier = session['metadata']['subscription_tier']

 # Update tenant subscription
 await update_tenant_subscription(tenant_id, new_tier, session)

 elif event['type'] == 'invoice.payment_failed':
 # Handle failed payment
 invoice = event['data']['object']
 await handle_payment_failure(invoice)

 return {"status": "success"}

```

## 10. Performance Optimization Strategies

### 10.1 Database Query Optimization

python

```

services/common/database_optimizations.py
from sqlalchemy import create_engine, event
from sqlalchemy.orm import sessionmaker
from sqlalchemy.pool import QueuePool
import logging

Connection pooling configuration for QNAP constraints
DATABASE_CONFIG = {
 "pool_size": 5, # Limit connections due to QNAP resources
 "max_overflow": 10, # Maximum additional connections
 "pool_pre_ping": True, # Verify connections before use
 "pool_recycle": 3600, # Recycle connections every hour
}

engine = create_engine(
 DATABASE_URL,
 **DATABASE_CONFIG,
 echo=False # Set to True for query debugging
)

Query performance monitoring
@event.listens_for(engine, "before_cursor_execute")
def receive_before_cursor_execute(conn, cursor, statement, parameters, context, executemany):
 context._query_start_time = time.time()

@event.listens_for(engine, "after_cursor_execute")
def receive_after_cursor_execute(conn, cursor, statement, parameters, context, executemany):
 total = time.time() - context._query_start_time
 if total > 1.0: # Log slow queries
 logging.warning(f"Slow query ({total:.2f}s): {statement[:200]}...")

Optimized tenant data access patterns
class TenantDataAccess:
 def __init__(self, tenant_id: str, db: Session):
 self.tenant_id = tenant_id
 self.db = db

 def get_dashboard_data(self) -> Dict[str, Any]:
 """Optimized dashboard data retrieval with single query"""

 # Single query to get all dashboard metrics
 result = self.db.execute(text("""
 WITH tenant_stats AS (

```

```

SELECT
 COUNT(DISTINCT o.id) as total_orders,
 COALESCE(SUM(o.total_amount), 0) as total_revenue,
 COUNT(DISTINCT p.id) as total_products,
 COUNT(DISTINCT d.id) as total_designs
FROM {schema}.orders o
FULL OUTER JOIN {schema}.products p ON true
FULL OUTER JOIN {schema}.designs d ON true
WHERE o.created_at >= CURRENT_DATE - INTERVAL '30 days'
),
recent_activity AS (
SELECT
 'order' as activity_type,
 id::text as activity_id,
 created_at,
 total_amount::text as activity_data
FROM {schema}.orders
WHERE created_at >= CURRENT_DATE - INTERVAL '7 days'

UNION ALL

SELECT
 'design' as activity_type,
 id::text as activity_id,
 created_at,
 name as activity_data
FROM {schema}.designs
WHERE created_at >= CURRENT_DATE - INTERVAL '7 days'
ORDER BY created_at DESC
LIMIT 10
)
SELECT
 json_build_object(
 'stats', row_to_json(ts),
 'recent_activity', json_agg(ra)
) as dashboard_data
FROM tenant_stats ts, recent_activity ra
GROUP BY ts.total_orders, ts.total_revenue, ts.total_products, ts.total_designs
""".format(schema=f"tenant_{self.tenant_id}"))

```

return result.fetchone()[0]

## 10.2 Caching Strategy (Redis)

python

```
services/common/cache.py
import redis
import json
import hashlib
from functools import wraps
from typing import Any, Optional, Callable
import pickle

redis_client = redis.Redis.from_url(REDIS_URL, decode_responses=False)

class TenantCache:
 def __init__(self, tenant_id: str):
 self.tenant_id = tenant_id
 self.key_prefix = f"tenant:{tenant_id}"

 def cache_key(self, key: str) -> str:
 """Generate tenant-scoped cache key"""
 return f"{self.key_prefix}:{key}"

 def get(self, key: str) -> Optional[Any]:
 """Get cached value"""
 cache_key = self.cache_key(key)
 data = redis_client.get(cache_key)
 if data:
 return pickle.loads(data)
 return None

 def set(self, key: str, value: Any, expire: int = 3600) -> None:
 """Set cached value with expiration"""
 cache_key = self.cache_key(key)
 redis_client.setex(cache_key, expire, pickle.dumps(value))

 def delete(self, key: str) -> None:
 """Delete cached value"""
 cache_key = self.cache_key(key)
 redis_client.delete(cache_key)

 def invalidate_pattern(self, pattern: str) -> None:
 """Invalidate all keys matching pattern"""
 pattern_key = self.cache_key(pattern)
 keys = redis_client.keys(pattern_key)
 if keys:
 redis_client.delete(*keys)
```

```
def cache_tenant_data(expire: int = 3600, key_suffix: str = None):
 """Decorator for caching tenant-specific data"""
 def decorator(func: Callable) -> Callable:
 @wraps(func)
 async def wrapper(tenant_id: str, *args, **kwargs):
 cache = TenantCache(tenant_id)

 # Generate cache key from function name and parameters
 if key_suffix:
 cache_key = f"{func.__name__}:{key_suffix}"
 else:
 # Create hash of parameters for unique key
 params_hash = hashlib.md5(
 json.dumps({"args": args, "kwargs": kwargs}, sort_keys=True).encode()
).hexdigest()[:8]
 cache_key = f"{func.__name__}:{params_hash}"

 # Try to get from cache
 cached_result = cache.get(cache_key)
 if cached_result is not None:
 return cached_result

 # Execute function and cache result
 result = await func(tenant_id, *args, **kwargs)
 cache.set(cache_key, result, expire)

 return result
 return wrapper
 return decorator

Usage examples
@cache_tenant_data(expire=1800, key_suffix="analytics")
async def get_tenant_analytics(tenant_id: str, date_range: str) -> Dict[str, Any]:
 """Cached analytics data"""
 # Expensive analytics computation here
 pass

@cache_tenant_data(expire=300) # 5 minutes for frequently changing data
async def get_etsy_listings(tenant_id: str, page: int = 1) -> Dict[str, Any]:
 """Cached Etsy listings with pagination"""
 # API call to Etsy
 pass
```

## **11. Deployment and Operations**

### **11.1 Production Deployment Script for QNAP**

```
bash
```

```
#!/bin/bash
deploy-production.sh

set -e

QNAP_HOST="your-qnap-ip"
QNAP_USER="admin"
PROJECT_PATH="/share/Container/etsy-saas-prod"
BACKUP_PATH="/share/Container/backups"

echo "🚀 Starting production deployment to QNAP NAS..."

Create deployment directory
ssh $QNAP_USER@$QNAP_HOST "mkdir -p $PROJECT_PATH"

Backup current deployment if exists
ssh $QNAP_USER@$QNAP_HOST "
if [-d $PROJECT_PATH/current]; then
 echo '📦 Creating backup of current deployment...'
 mkdir -p $BACKUP_PATH
 tar -czf $BACKUP_PATH/backup-$(date +%Y%m%d-%H%M%S).tar.gz -C $PROJECT_PATH current
fi
"

Upload new deployment files
echo "📦 Uploading deployment files..."
rsync -avz --exclude='\.git' --exclude='node_modules' --exclude='__pycache__' \
./ $QNAP_USER@$QNAP_HOST:$PROJECT_PATH/new/

Switch to new deployment
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH
if [-d current]; then
 mv current previous
fi
mv new current
"

Start services
echo "⌚ Starting production services..."
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH/current
docker-compose -f docker-compose.prod.yml down || true
```

```

docker-compose -f docker-compose.prod.yml pull
docker-compose -f docker-compose.prod.yml up -d
"

Health check
echo "🏥 Running health checks..."
sleep 30
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH/current
./scripts/health-check.sh
"

Cleanup old images
ssh $QNAP_USER@$QNAP_HOST "
echo '🧹 Cleaning up old Docker images...'
docker system prune -f
docker image prune -f
"

echo "✅ Production deployment completed successfully!"
echo "🌐 Services available at: https://your-domain.com"
echo "💻 Monitor logs with: ssh $QNAP_USER@$QNAP_HOST 'cd $PROJECT_PATH/current && docker-compos
 CONSTRAINT valid_tier CHECK (subscription_tier IN ('starter', 'professional', 'enterprise'))
);
```

```

CREATE TABLE IF NOT EXISTS core.tenant_users (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id) ON DELETE CASCADE,
 email VARCHAR(255) NOT NULL,
 hashed_password VARCHAR(255) NOT NULL,
 first_name VARCHAR(100),
 last_name VARCHAR(100),
 role VARCHAR(50) DEFAULT 'user',
 permissions JSONB DEFAULT '[{"user:read"}]',
 is_active BOOLEAN DEFAULT TRUE,
 email_verified BOOLEAN DEFAULT FALSE,
 last_login TIMESTAMP,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 UNIQUE(tenant_id, email),
 CONSTRAINT valid_role CHECK (role IN ('admin', 'user', 'viewer'))
);
```

```
-- Subscription tracking
CREATE TABLE IF NOT EXISTS core.subscriptions (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id) ON DELETE CASCADE,
 stripe_subscription_id VARCHAR(100) UNIQUE,
 status VARCHAR(50) NOT NULL,
 tier VARCHAR(50) NOT NULL,
 current_period_start TIMESTAMP,
 current_period_end TIMESTAMP,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Usage tracking
CREATE TABLE IF NOT EXISTS core.tenant_usage (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id) ON DELETE CASCADE,
 metric_name VARCHAR(100) NOT NULL,
 metric_value INTEGER NOT NULL,
 recorded_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 period_start TIMESTAMP NOT NULL,
 period_end TIMESTAMP NOT NULL,
 UNIQUE(tenant_id, metric_name, period_start)
);

-- Audit log
CREATE TABLE IF NOT EXISTS core.audit_log (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id) ON DELETE CASCADE,
 user_id UUID,
 action VARCHAR(100) NOT NULL,
 resource_type VARCHAR(100),
 resource_id VARCHAR(100),
 details JSONB,
 ip_address INET,
 user_agent TEXT,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Create indexes
CREATE INDEX IF NOT EXISTS idx_tenants_subdomain ON core.tenants(subdomain);
CREATE INDEX IF NOT EXISTS idx_tenant_users_tenant_email ON core.tenant_users(tenant_id, email);
CREATE INDEX IF NOT EXISTS idx_subscriptions_tenant ON core.subscriptions(tenant_id);
```

```
CREATE INDEX IF NOT EXISTS idx_usage_tenant_metric ON core.tenant_usage(tenant_id, metric_name, period_|
CREATE INDEX IF NOT EXISTS idx_audit_tenant_created ON core.audit_log(tenant_id, created_at);
```

```
-- Create function to automatically update updated_at timestamps
```

```
CREATE OR REPLACE FUNCTION update_updated_at_column()
```

```
RETURNS TRIGGER AS \$\$
```

```
BEGIN
```

```
 NEW.updated_at = CURRENT_TIMESTAMP;
```

```
 RETURN NEW;
```

```
END;
```

```
\$\$ language 'plpgsql';
```

```
-- Create triggers for updated_at
```

```
CREATE TRIGGER update_tenants_updated_at BEFORE UPDATE ON core.tenants
```

```
 FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_tenant_users_updated_at BEFORE UPDATE ON core.tenant_users
```

```
 FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
CREATE TRIGGER update_subscriptions_updated_at BEFORE UPDATE ON core.subscriptions
```

```
 FOR EACH ROW EXECUTE FUNCTION update_updated_at_column();
```

```
EOF
```

```
 log "✅ Database initialization scripts created"
```

```
}
```

```
Start services
```

```
start_services() {
```

```
 log "🚀 Starting production services..."
```

```
Pull latest images
```

```
info "Pulling latest container images..."
```

```
docker-compose -f "$COMPOSE_FILE" --env-file "$ENV_FILE" pull
```

```
Start infrastructure services first
```

```
info "Starting infrastructure services..."
```

```
docker-compose -f "$COMPOSE_FILE" --env-file "$ENV_FILE" up -d postgres redis minio
```

```
Wait for infrastructure to be ready
```

```
info "Waiting for infrastructure services to be ready..."
```

```
sleep 30
```

```
Run database migrations
```

```
info "Running database migrations..."
```

```

Add migration commands here if needed

Start application services
info "Starting application services..."
docker-compose -f "$COMPOSE_FILE" --env-file "$ENV_FILE" up -d

log "✅ All services started successfully"
}

Health checks
run_health_checks() {
 log "📝 Running health checks..."

 services=("api-gateway" "auth-service" "etsy-service" "design-service" "analytics-service" "postgres" "redis"
max_retries=10
retry_interval=15

 for service in "${services[@]}"; do
 info "Checking health of $service..."

 retries=0
 while [[$retries -lt $max_retries]]; do
 if docker-compose -f "$COMPOSE_FILE" ps | grep -q "$service.*Up.*healthy\|$service.*Up"; then
 log "✅ $service is healthy"
 break
 fi

 retries=$((retries + 1))
 if [[$retries -eq $max_retries]]; then
 error "🔴 $service failed health check after $max_retries attempts"
 return 1
 fi

 warning "⏳ $service not ready, retrying in ${retry_interval}s... (attempt ${retries}/${max_retries})"
 sleep $retry_interval
 done
 done

 log "✅ All health checks passed"
}

Setup monitoring dashboards
setup_dashboards() {
 log "📊 Setting up monitoring dashboards..."
}

```

```
Create basic dashboard for tenant metrics
cat > ./monitoring/grafana/dashboards/tenant-overview.json << 'EOF'
{
 "dashboard": {
 "id": null,
 "title": "Etsy SaaS - Tenant Overview",
 "tags": ["etsy", "saas"],
 "timezone": "browser",
 "panels": [
 {
 "id": 1,
 "title": "Active Tenants",
 "type": "stat",
 "targets": [
 {
 "expr": "count(tenant_active{status=\"active\"})",
 "legendFormat": "Active Tenants"
 }
]
 },
 {
 "id": 2,
 "title": "API Requests per Minute",
 "type": "graph",
 "targets": [
 {
 "expr": "rate(http_requests_total[1m])",
 "legendFormat": "{{tenant_id}}"
 }
]
 },
 {
 "id": 3,
 "title": "Database Connections",
 "type": "graph",
 "targets": [
 {
 "expr": "pg_stat_activity_count",
 "legendFormat": "Connections"
 }
]
 }
],
 }
}
```

```
"time": {
 "from": "now-1h",
 "to": "now"
},
"refresh": "30s"
}
}
EOF

log "✅ Monitoring dashboards created"
}

Main execution
main() {
 log "⭐ Starting Etsy Seller Automater SaaS Production Deployment"

 check_qnap_environment
 pre_flight_checks
 setup_directories
 setup_monitoring_config
 init_database
 start_services
 run_health_checks
 setup_dashboards

 log "🎉 Production deployment completed successfully!"

 echo ""
 info "📋 Service URLs:"
 info " • Main Application: https://yourdomain.com"
 info " • API Gateway: https://yourdomain.com/api"
 info " • Traefik Dashboard: http://your-qnap-ip:8080"
 info " • Grafana Monitoring: https://monitoring.yourdomain.com"
 info " • Prometheus Metrics: https://prometheus.yourdomain.com"
 info " • MinIO Console: https://minio.yourdomain.com"

 echo ""
 info "📊 Management Commands:"
 info " • View logs: docker-compose -f $COMPOSE_FILE logs -f [service]"
 info " • Scale service: docker-compose -f $COMPOSE_FILE up -d --scale service-name=3"
 info " • Update services: docker-compose -f $COMPOSE_FILE pull && docker-compose -f $COMPOSE_FILE u
 info " • Stop all: docker-compose -f $COMPOSE_FILE down"

 echo ""
```

```

warning "🔒 Security Reminders:"
warning " • Change all default passwords in $ENV_FILE"
warning " • Configure firewall rules for production"
warning " • Set up SSL certificates for all domains"
warning " • Enable backup automation"
warning " • Configure log rotation"
}

Run main function
main "$@"

```

## Conclusion

This comprehensive technical deep dive provides a complete multi-tenant SaaS architecture for your Etsy Seller Automater platform using Python and Go microservices. The architecture is specifically optimized for deployment on your QNAP TBS-464 NAS while maintaining scalability and professional-grade features.

### Key Highlights:

1. **Multi-Tenant Architecture:** Schema-per-tenant approach for optimal data isolation
2. **Microservices Design:** Modular services in Python (FastAPI) and Go for different concerns
3. **QNAP Optimization:** Resource constraints considered with memory/CPU limits
4. **Security First:** JWT authentication, rate limiting, tenant isolation
5. **Production Ready:** Monitoring, logging, health checks, automated deployment
6. **Scalable:** Horizontal scaling strategy and load balancing
7. **Cost Effective:** Leverages existing QNAP hardware with cloud migration path

The architecture supports your current features while adding enterprise-level capabilities like billing, analytics, job processing, and comprehensive monitoring. The deployment scripts and configuration files provide a complete production-ready solution that can start small on your QNAP NAS and scale to cloud infrastructure as your SaaS grows.

network  
restart: unless-stopped

## API Gateway

```

api-gateway: build: context: ./services/gateway dockerfile: Dockerfile.prod container_name: etsy-
gateway environment: - REDIS_URL=redis://redis:6379/0 -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
JWT_SECRET=${JWT_SECRET} deploy: resources: limits: memory: 512M cpus: '1.0' reservations:

```

```
memory: 256M cpus: '0.5' labels: - "traefik.enable=true" - "traefik.http.routers.api-gateway.rule=PathPrefix(/api)" - "traefik.http.routers.api-gateway.tls=true" - "traefik.http.routers.api-gateway.tls.certresolver=letsencrypt" depends_on: - postgres - redis networks: - etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f", "http://localhost:8080/health"] interval: 30s timeout: 10s retries: 3
```

## Authentication Service

```
auth-service: build: context: ./services/auth dockerfile: Dockerfile.prod container_name: etsy-auth environment: -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
REDIS_URL=redis://redis:6379/1 - JWT_SECRET=${JWT_SECRET} - BCRYPT_ROUNDS=12 deploy:
resources: limits: memory: 1G cpus: '1.0' reservations: memory: 512M cpus: '0.5' depends_on: -
postgres - redis networks: - etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f", "http://localhost:8001/health"] interval: 30s timeout: 10s retries: 3
```

## Etsy Integration Service

```
etsy-service: build: context: ./services/etsy dockerfile: Dockerfile.prod container_name: etsy-integration environment: -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
REDIS_URL=redis://redis:6379/2 - ETSY_API_BASE_URL=https://api.etsy.com/v3 -
ETSY_RATE_LIMIT_CALLS=10000 - ETSY_RATE_LIMIT_WINDOW=3600 deploy: resources: limits:
memory: 1G cpus: '1.0' reservations: memory: 512M cpus: '0.5' depends_on: - postgres - redis
networks: - etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f",
"http://localhost:8002/health"] interval: 30s timeout: 10s retries: 3
```

## Design Management Service

```
design-service: build: context: ./services/design dockerfile: Dockerfile.prod container_name: etsy-design environment: -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
MINIO_ENDPOINT=minio:9000 - MINIO_ACCESS_KEY=${MINIO_ACCESS_KEY} -
MINIO_SECRET_KEY=${MINIO_SECRET_KEY} - MINIO_SECURE=false deploy: resources: limits:
memory: 2G cpus: '2.0' reservations: memory: 1G cpus: '1.0' depends_on: - postgres - minio networks:
- etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f",
"http://localhost:8003/health"] interval: 30s timeout: 10s retries: 3
```

## Analytics Service

```
analytics-service: build: context: ./services/analytics dockerfile: Dockerfile.prod container_name: etsy-analytics environment: -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
REDIS_URL=redis://redis:6379/3 deploy: resources: limits: memory: 1G cpus: '1.0' reservations: memory: 512M cpus: '0.5'
depends_on: - postgres - redis networks: - etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f", "http://localhost:8004/health"] interval: 30s timeout:
10s retries: 3
```

## Job Processing Workers

```
job-worker: build: context: ./services/jobs dockerfile: Dockerfile.prod container_name: etsy-worker environment: -
DATABASE_URL=postgresql://postgres:${POSTGRES_PASSWORD}@postgres:5432/etsy_saas -
REDIS_URL=redis://redis:6379/4 - WORKER_CONCURRENCY=4 - WORKER_MAX_MEMORY=1G
deploy: replicas: 2 resources: limits: memory: 1G cpus: '1.0' reservations: memory: 512M cpus: '0.5'
depends_on: - redis - postgres networks: - etsy-network restart: unless-stopped
```

## Notification Service

```
notification-service: build: context: ./services/notifications dockerfile: Dockerfile.prod container_name: etsy-notifications environment: -
REDIS_URL=redis://redis:6379/5 deploy: resources: limits: memory: 512M cpus: '0.5' reservations: memory: 256M cpus: '0.25'
depends_on: - redis networks: - etsy-network restart: unless-stopped healthcheck: test: ["CMD", "curl", "-f", "http://localhost:8005/health"] interval: 30s timeout:
10s retries: 3
```

## Infrastructure Services

```
postgres: image: postgres:15 container_name: etsy-postgres environment: -
POSTGRES_DB=etsy_saas - POSTGRES_USER=postgres -
POSTGRES_PASSWORD=${POSTGRES_PASSWORD} - POSTGRES_SHARED_BUFFERS=512MB -
POSTGRES_EFFECTIVE_CACHE_SIZE=2GB - POSTGRES_MAINTENANCE_WORK_MEM=128MB -
POSTGRES_WAL_BUFFERS=16MB volumes: - postgres_data:/var/lib/postgresql/data - ./init-scripts:/docker-entrypoint-initdb.d - ./postgres-config:/etc/postgresql/postgresql.conf ports: -
"5432:5432" deploy: resources: limits: memory: 3G cpus: '2.0' reservations: memory: 2G cpus: '1.0'
networks: - etsy-# Multi-Tenant Etsy Seller Automater Backend Architecture Deep Dive
```

## Executive Summary

This document outlines the backend architecture for transforming the existing Etsy Seller Automater into a scalable multi-tenant SaaS platform, leveraging Python and Go microservices deployed on

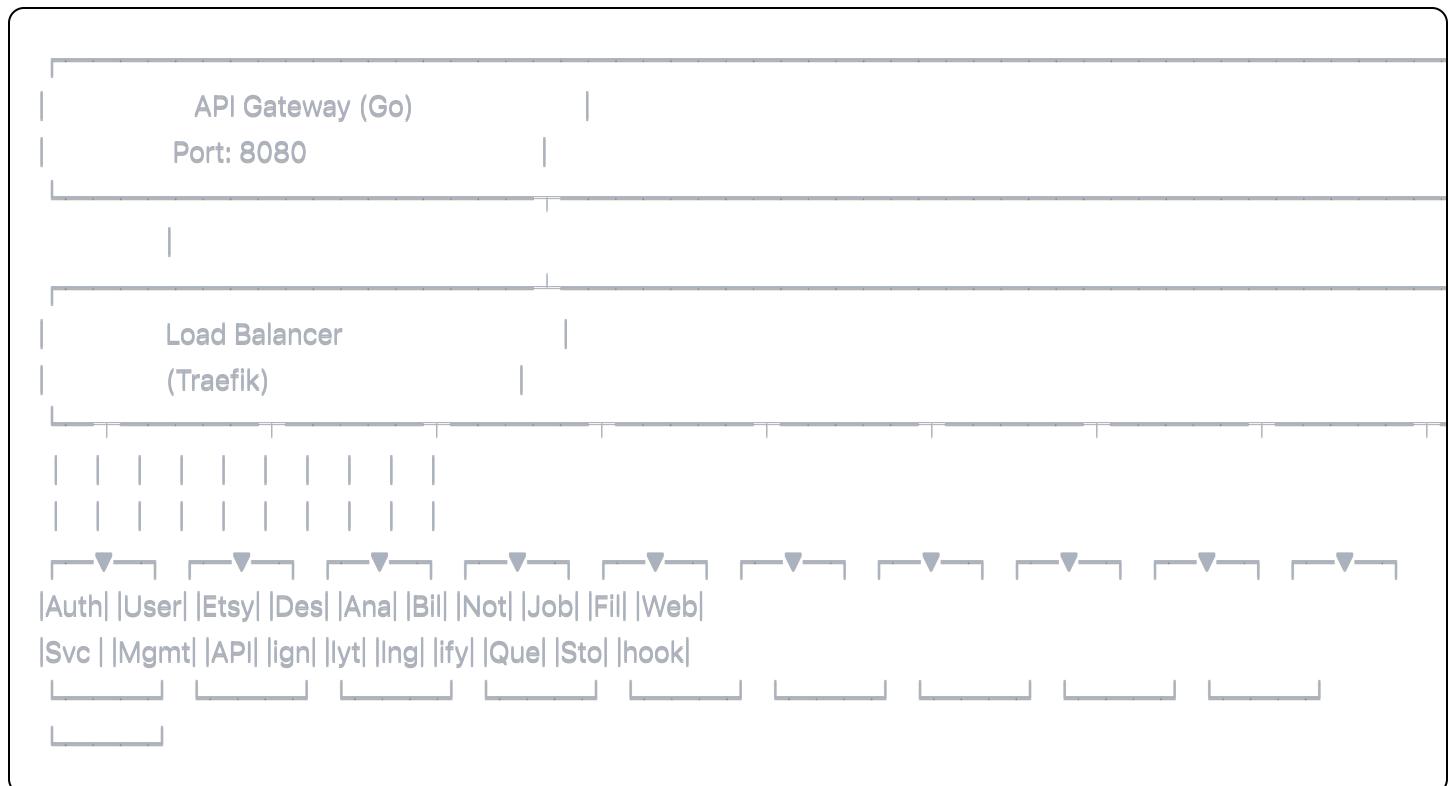
## Current State Analysis

### Existing Architecture

- **Framework:** FastAPI (Python)
- **Database:** PostgreSQL
- **Authentication:** OAuth 2.0 with PKCE for Etsy API
- **Frontend:** React with Tailwind CSS
- **Deployment:** Docker Compose
- **Key Features:**
  - Etsy shop analytics
  - Design management and mockup creation
  - Listing automation
  - Mask creator for product images

### Proposed Multi-Tenant Architecture

#### 1. Service Architecture Overview



#### 2. Core Services Architecture

## 2.1 API Gateway Service (Go)

**Port:** 8080

**Responsibilities:** Request routing, rate limiting, authentication middleware, tenant isolation

go

```
// pkg/gateway/main.go
package main

import (
 "context"
 "log"
 "net/http"
 "time"

 "github.com/gin-gonic/gin"
 "github.com/redis/go-redis/v9"
)

type Gateway struct {
 router *gin.Engine
 redisClient *redis.Client
 services map[string]string
}

type TenantMiddleware struct {
 redis *redis.Client
}

func (tm *TenantMiddleware) ExtractTenant() gin.HandlerFunc {
 return func(c *gin.Context) {
 host := c.Request.Host
 subdomain := extractSubdomain(host)

 // Validate tenant exists
 tenantID, err := tm.redis.Get(c.Request.Context(),
 fmt.Sprintf("tenant:%s", subdomain)).Result()
 if err != nil {
 c.JSON(404, gin.H{"error": "Tenant not found"})
 c.Abort()
 return
 }

 c.Set("tenant_id", tenantID)
 c.Set("subdomain", subdomain)
 c.Next()
 }
}
```

```
func (g *Gateway) setupRoutes() {
 v1 := g.router.Group("/api/v1")
 v1.Use(g.tenantMiddleware.ExtractTenant())

 // Route to microservices
 v1.Any("/auth/*path", g.proxyToService("auth-service"))
 v1.Any("/etsy/*path", g.proxyToService("etsy-service"))
 v1.Any("/designs/*path", g.proxyToService("design-service"))
 v1.Any("/analytics/*path", g.proxyToService("analytics-service"))
}
```

## 2.2 Authentication & User Management Service (Python/FastAPI)

**Port:** 8001

**Database:** `tenant_users` schema in PostgreSQL

python

```
services/auth/main.py
from fastapi import FastAPI, Depends, HTTPException
from sqlalchemy.orm import Session
import jwt
from datetime import datetime, timedelta

app = FastAPI(title="Authentication Service")

class AuthService:
 def __init__(self, db: Session):
 self.db = db

 async def authenticate_tenant_user(self,
 tenant_id: str,
 email: str,
 password: str) -> dict:
 """Authenticate user within specific tenant context"""
 user = self.db.query(TenantUser).filter(
 TenantUser.tenant_id == tenant_id,
 TenantUser.email == email,
 TenantUser.is_active == True
).first()

 if not user or not verify_password(password, user.hashed_password):
 raise HTTPException(401, "Invalid credentials")

 # Generate tenant-scoped JWT
 payload = {
 "user_id": user.id,
 "tenant_id": tenant_id,
 "permissions": user.permissions,
 "exp": datetime.utcnow() + timedelta(hours=24)
 }

 token = jwt.encode(payload, JWT_SECRET, algorithm="HS256")

 return {
 "access_token": token,
 "token_type": "bearer",
 "user": user.to_dict(),
 "tenant": user.tenant.to_dict()
 }
```

```

Database Models
class Tenant(Base):
 __tablename__ = "tenants"

 id = Column(UUID(as_uuid=True), primary_key=True, default=uuid.uuid4)
 subdomain = Column(String(63), unique=True, nullable=False)
 company_name = Column(String(255), nullable=False)
 subscription_tier = Column(String(50), default="basic")
 created_at = Column(DateTime, default=datetime.utcnow)

 # Etsy Integration
 etsy_shop_id = Column(String(100))
 etsy_access_token = Column(Text)
 etsy_refresh_token = Column(Text)

class TenantUser(Base):
 __tablename__ = "tenant_users"

 id = Column(UUID(as_uuid=True), primary_key=True, default=uuid.uuid4)
 tenant_id = Column(UUID(as_uuid=True), ForeignKey("tenants.id"))
 email = Column(String(255), nullable=False)
 hashed_password = Column(String(255), nullable=False)
 permissions = Column(JSON) # Role-based permissions
 is_active = Column(Boolean, default=True)

```

## 2.3 Etsy Integration Service (Python/FastAPI)

**Port:** 8002

**Focus:** OAuth flows, API calls, rate limiting per tenant

python

```
services/etsy/main.py
import asyncio
import aiohttp
from fastapi import FastAPI, BackgroundTasks
from rq import Queue
import redis

app = FastAPI(title="Etsy Integration Service")

class EtsyService:
 def __init__(self, tenant_id: str):
 self.tenant_id = tenant_id
 self.base_url = "https://api.etsy.com/v3"

 @async def get_shop_analytics(self, date_range: str) -> dict:
 """Get shop analytics for tenant with rate limiting"""

 # Implement tenant-specific rate limiting
 rate_limit_key = f"etsy_rate_limit:{self.tenant_id}"

 async with aiohttp.ClientSession() as session:
 headers = await self._get_auth_headers()

 # Get receipt data
 receipts = await self._fetch_receipts(session, headers, date_range)

 # Process analytics
 analytics = await self._process_analytics_data(receipts)

 return {
 "tenant_id": self.tenant_id,
 "date_range": date_range,
 "total_revenue": analytics["revenue"],
 "total_orders": analytics["orders"],
 "top_products": analytics["top_products"],
 "monthly_breakdown": analytics["monthly"]
 }

 @async def create_listing(self, listing_data: dict) -> dict:
 """Create Etsy listing with tenant isolation"""

 # Validate listing data against tenant's subscription limits
 await self._validate_listing_limits()
```

```

Create listing via Etsy API
async with aiohttp.ClientSession() as session:
 headers = await self._get_auth_headers()

 response = await session.post(
 f"{self.base_url}/application/shops/{self.shop_id}/listings",
 headers=headers,
 json=listing_data
)

 result = await response.json()

Store listing in tenant database
await self._store_listing_record(result)

return result

Background job for bulk operations
@app.post("/bulk-listing-creation")
async def bulk_create_listings(
 listings: List[dict],
 tenant_id: str,
 background_tasks: BackgroundTasks
):
 """Queue bulk listing creation as background job"""

 redis_conn = redis.Redis(host='redis', port=6379, db=0)
 queue = Queue('bulk_operations', connection=redis_conn)

 job = queue.enqueue(
 'tasks.bulk_create_listings',
 listings,
 tenant_id,
 job_timeout='30m'
)

 return {"job_id": job.id, "status": "queued"}

```

## 2.4 Design Management Service (Go)

**Port:** 8003

**Focus:** File storage, image processing, mockup generation

go

```
// services/design/main.go
package main

import (
 "context"
 "fmt"
 "io"
 "path/filepath"

 "github.com/gin-gonic/gin"
 "github.com/minio/minio-go/v7"
 "gorm.io/gorm"
)

type DesignService struct {
 db *gorm.DB
 minioClient *minio.Client
 router *gin.Engine
}

type Design struct {
 ID uint `gorm:"primaryKey"`
 TenantID string `gorm:"index"`
 Name string
 FilePath string
 FileSize int64
 ContentType string
 Tags []string `gorm:"serializer:json"`
 CreatedAt time.Time
}

func (ds *DesignService) uploadDesign(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 file, header, err := c.Request.FormFile("design")
 if err != nil {
 c.JSON(400, gin.H{"error": "Invalid file"})
 return
 }
 defer file.Close()

 // Generate tenant-scoped file path
 fileName := fmt.Sprintf("%s/%s/%s",
 tenantID,
 header.Filename,
 file.Size)
}
```

```

tenantID, "designs", header.Filename)

// Upload to MinIO
_, err = ds.minioClient.PutObject(
 context.Background(),
 "tenant-designs",
 fileName,
 file,
 header.Size,
 minio.PutObjectOptions{
 ContentType: header.Header.Get("Content-Type"),
 },
)
}

if err != nil {
 c.JSON(500, gin.H{"error": "Upload failed"})
 return
}

// Save metadata to database
design := Design{
 TenantID: tenantID,
 Name: header.Filename,
 FilePath: fileName,
 FileSize: header.Size,
 ContentType: header.Header.Get("Content-Type"),
}
ds.db.Create(&design)

c.JSON(201, design)
}

func (ds *DesignService) generateMockup(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
 designID := c.Param("id")

 in Get design from database
 var design Design
 result := ds.db.Where("tenant_id = ? AND id = ?",
 tenantID, designID).First(&design)

 if result.Error != nil {
 c.JSON(404, gin.H{"error": "Design not found"})
 }
}

```

```

 return
}

// Queue mockup generation
job := MockupGenerationJob{
 TenantID: tenantID,
 DesignID: designID,
 FilePath: design.FilePath,
}

// Submit to job queue (Redis/RQ)
jobID, err := ds.queueMockupJob(job)
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to queue job"})
 return
}

c.JSON(202, gin.H{
 "job_id": jobID,
 "status": "queued",
 "message": "Mockup generation started"
})
}

```

## 2.5 Analytics Service (Python/FastAPI)

**Port:** 8004

**Focus:** Data aggregation, reporting, business intelligence

python



```
analytics = {
 "summary": {
 "total_revenue": df['daily_revenue'].sum(),
 "total_orders": df['daily_orders'].sum(),
 "avg_daily_revenue": df['daily_revenue'].mean(),
 "revenue_growth": self._calculate_growth_rate(df)
 },
 "daily_breakdown": df.to_dict('records'),
 "top_performing_days": df.nlargest(5, 'daily_revenue').to_dict('records')
}

return analytics
```

```
async def get_product_performance(self) -> dict:
 """Analyze product performance metrics"""

 query = text("""
```

```
SELECT
 p.name,
 p.etsy_listing_id,
 COUNT(oi.id) as total_sales,
 SUM(oi.quantity) as units_sold,
 SUM(oi.price * oi.quantity) as total_revenue,
 AVG(oi.price) as avg_price
FROM products p
LEFT JOIN order_items oi ON p.id = oi.product_id
LEFT JOIN tenant_orders o ON oi.order_id = o.id
WHERE o.tenant_id = :tenant_id
GROUP BY p.id, p.name, p.etsy_listing_id
ORDER BY total_revenue DESC
""")
```

```
result = self.db.execute(query, {"tenant_id": self.tenant_id})
```

```
products = []
for row in result:
 products.append({
 "name": row.name,
 "etsy_listing_id": row.etsy_listing_id,
 "total_sales": row.total_sales,
 "units_sold": row.units_sold,
 "total_revenue": float(row.total_revenue or 0),
 "avg_price": float(row.avg_price or 0)
 })
```

```

return {
 "top_products": products[:10],
 "total_products": len(products),
 "products_with_sales": len([p for p in products if p['total_sales'] > 0])
}

@app.get("/analytics/dashboard")
async def get_dashboard_data(
 tenant_id: str = Depends(get_tenant_id),
 date_range: str = "30d"
):
 """Get comprehensive dashboard analytics"""

 analytics_service = AnalyticsService(tenant_id, get_db())

 # Calculate date range
 end_date = datetime.utcnow()
 days = int(date_range.replace('d', ''))
 start_date = end_date - timedelta(days=days)

 sales_data = await analytics_service.get_sales_analytics(start_date, end_date)
 product_data = await analytics_service.get_product_performance()

 return {
 "tenant_id": tenant_id,
 "date_range": date_range,
 "sales": sales_data,
 "products": product_data,
 "generated_at": datetime.utcnow().isoformat()
}

```

### 3. Database Architecture

#### 3.1 Multi-Tenant Schema Design (PostgreSQL)

sql

```

-- Core tenant management
CREATE SCHEMA IF NOT EXISTS core;

CREATE TABLE core.tenants (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 subdomain VARCHAR(63) UNIQUE NOT NULL,
 company_name VARCHAR(255) NOT NULL,
 subscription_tier VARCHAR(50) DEFAULT 'basic',
 database_schema VARCHAR(63) NOT NULL,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
 updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

CREATE TABLE core.tenant_users (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 tenant_id UUID REFERENCES core.tenants(id),
 email VARCHAR(255) NOT NULL,
 hashed_password VARCHAR(255) NOT NULL,
 role VARCHAR(50) DEFAULT 'user',
 permissions JSONB,
 is_active BOOLEAN DEFAULT TRUE,
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Function to create tenant schema
CREATE OR REPLACE FUNCTION create_tenant_schema(tenant_id UUID, schema_name VARCHAR)
RETURNS VOID AS $$
BEGIN
 -- Create schema
 EXECUTE format('CREATE SCHEMA IF NOT EXISTS %I', schema_name);

 -- Create tenant-specific tables
 EXECUTE format(
 CREATE TABLE %I.orders (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 etsy_receipt_id BIGINT,
 total_amount DECIMAL(10,2),
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
), schema_name);

 EXECUTE format(
 CREATE TABLE %I.products (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```
 name VARCHAR(255) NOT NULL,
 etsy_listing_id BIGINT,
 price DECIMAL(10,2),
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)', schema_name);
```

```
EXECUTE format('
```

```
CREATE TABLE %I.designs (
 id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
 name VARCHAR(255) NOT NULL,
 file_path VARCHAR(500),
 file_size BIGINT,
 tags TEXT[],
 created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
)', schema_name);
```

```
END;
```

```
$$ LANGUAGE plpgsql;
```

## 4. Container Orchestration (Docker Compose)

```
yaml
```

```
docker-compose.yml
version: '3.8'

services:
API Gateway
api-gateway:
build:
 context: ./services/gateway
 dockerfile: Dockerfile
ports:
- "8080:8080"
environment:
- REDIS_URL=redis://redis:6379
depends_on:
- redis
networks:
- etsy-network

Authentication Service
auth-service:
build:
 context: ./services/auth
 dockerfile: Dockerfile
ports:
- "8001:8001"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- JWT_SECRET=your-jwt-secret-here
- REDIS_URL=redis://redis:6379
depends_on:
- postgres
- redis
networks:
- etsy-network

Etsy Integration Service
etsy-service:
build:
 context: ./services/etsy
 dockerfile: Dockerfile
ports:
- "8002:8002"
environment:
```

```
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- REDIS_URL=redis://redis:6379
- ETSY_API_BASE_URL=https://api.etsy.com/v3

depends_on:
- postgres
- redis

networks:
- etsy-network

Design Management Service (Go)
design-service:
build:
context: ./services/design
dockerfile: Dockerfile
ports:
- "8003:8003"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
- MINIO_ENDPOINT=minio:9000
- MINIO_ACCESS_KEY=minioadmin
- MINIO_SECRET_KEY=minioadmin
depends_on:
- postgres
- minio
networks:
- etsy-network

Analytics Service
analytics-service:
build:
context: ./services/analytics
dockerfile: Dockerfile
ports:
- "8004:8004"
environment:
- DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
depends_on:
- postgres
networks:
- etsy-network

Job Queue Worker
job-worker:
build:
```

```
context: ./services/jobs
dockerfile: Dockerfile
environment:
 - REDIS_URL=redis://redis:6379
 - DATABASE_URL=postgresql://postgres:password@postgres:5432/etsy_saas
depends_on:
 - redis
 - postgres
networks:
 - etsy-network

Infrastructure Services

postgres:
 image: postgres:15
 environment:
 - POSTGRES_DB=etsy_saas
 - POSTGRES_USER=postgres
 - POSTGRES_PASSWORD=password
 volumes:
 - postgres_data:/var/lib/postgresql/data
 - ./init.sql:/docker-entrypoint-initdb.d/init.sql
 ports:
 - "5432:5432"
 networks:
 - etsy-network

redis:
 image: redis:7-alpine
 ports:
 - "6379:6379"
 volumes:
 - redis_data:/data
 networks:
 - etsy-network

minio:
 image: minio/minio
 ports:
 - "9000:9000"
 - "9001:9001"
 volumes:
 - minio_data:/data
 environment:
 - MINIO_ROOT_USER=minioadmin
```

```
- MINIO_ROOT_PASSWORD=minioadmin
command: server /data --console-address ":9001"
networks:
- etsy-network

Reverse Proxy
traefik:
image: traefik:v2.10
command:
- --providers.docker=true
- --providers.docker.exposedbydefault=false
- --entrypoints.web.address=:80
- --entrypoints.websecure.address=:443
ports:
- "80:80"
- "443:443"
- "8080:8080" # Dashboard
volumes:
- /var/run/docker.sock:/var/run/docker.sock:ro
networks:
- etsy-network

volumes:
postgres_data:
redis_data:
minio_data:

networks:
etsy-network:
driver: bridge
```

## 5. Deployment on QNAP NAS

### 5.1 QNAP-Specific Configuration

```
bash
```

```
#!/bin/bash
deploy-qnap.sh

Install Container Station if not already installed
Access QNAP Web UI -> App Center -> Container Station

Create project directory
ssh admin@qnap-ip "mkdir -p /share/Container/etsy-saas"

Copy docker-compose files
scp docker-compose.yml admin@qnap-ip:/share/Container/etsy-saas/
scp docker-compose.override.yml admin@qnap-ip:/share/Container/etsy-saas/

Create persistent volumes on QNAP
ssh admin@qnap-ip "mkdir -p /share/Container/volumes/{postgres,redis,minio}"

Set proper permissions
ssh admin@qnap-ip "chmod -R 755 /share/Container/etsy-saas"

Start services
ssh admin@qnap-ip "cd /share/Container/etsy-saas && docker-compose up -d"
```

## 5.2 QNAP Resource Optimization

yaml

```
docker-compose.override.yml - QNAP specific optimizations
version: '3.8'

services:
 api-gateway:
 deploy:
 resources:
 limits:
 memory: 256M
 cpus: '0.5'

 auth-service:
 deploy:
 resources:
 limits:
 memory: 512M
 cpus: '0.5'

 etsy-service:
 deploy:
 resources:
 limits:
 memory: 512M
 cpus: '0.5'

 design-service:
 deploy:
 resources:
 limits:
 memory: 1G
 cpus: '1.0'

 postgres:
 deploy:
 resources:
 limits:
 memory: 2G
 cpus: '1.0'

 environment:
 - POSTGRES_SHARED_BUFFERS=256MB
 - POSTGRES_EFFECTIVE_CACHE_SIZE=1GB
```

## 6. Monitoring and Observability

### 6.1 Health Checks

```
python
```

```
services/common/health.py
from fastapi import FastAPI
import asyncio
import psycopg2
import redis

class HealthCheck:
 def __init__(self, app: FastAPI):
 self.app = app
 self.setup_routes()

 def setup_routes(self):
 @self.app.get("/health")
 async def health_check():
 checks = {
 "database": await self.check_database(),
 "redis": await self.check_redis(),
 "service": "healthy"
 }

 all_healthy = all(
 status == "healthy"
 for status in checks.values()
)

 return {
 "status": "healthy" if all_healthy else "unhealthy",
 "checks": checks
 }

 async def check_database(self) -> str:
 try:
 conn = psycopg2.connect(DATABASE_URL)
 conn.close()
 return "healthy"
 except Exception as e:
 return f"unhealthy: {str(e)}"

 async def check_redis(self) -> str:
 try:
 r = redis.Redis.from_url(REDIS_URL)
 r.ping()
 return "healthy"
```

```
except Exception as e:
 return f"unhealthy: {str(e)}"
```

## 7. Security Considerations

### 7.1 Tenant Isolation Security

python

```
services/common/security.py
from functools import wraps
import jwt
from fastapi import HTTPException

def require_tenant_access(required_permissions: List[str] = None):
 def decorator(func):
 @wraps(func)
 async def wrapper(*args, **kwargs):
 # Extract JWT token
 token = extract_token_from_request()

 try:
 payload = jwt.decode(token, JWT_SECRET, algorithms=["HS256"])
 tenant_id = payload.get("tenant_id")
 user_permissions = payload.get("permissions", [])

 # Verify tenant access
 if not tenant_id:
 raise HTTPException(403, "No tenant access")

 # Check specific permissions if required
 if required_permissions:
 if not any(perm in user_permissions for perm in required_permissions):
 raise HTTPException(403, "Insufficient permissions")

 # Add to request context
 kwargs['tenant_id'] = tenant_id
 kwargs['user_permissions'] = user_permissions

 except jwt.InvalidTokenError:
 raise HTTPException(401, "Invalid token")

 return await func(*args, **kwargs)
 return wrapper
 return decorator

Usage example
@app.post("/designs/upload")
@require_tenant_access(["design:create"])
async def upload_design(tenant_id: str, user_permissions: List[str], file: UploadFile = File(...)):
 """Upload design with tenant isolation"""

```

```
Process upload with tenant context
```

```
pass
```

## 8. Background Job Processing Architecture

### 8.1 Job Queue System (Python/RQ + Redis)

```
python
```

```
services/jobs/worker.py
from rq import Worker, Queue, Connection
import redis
from typing import Dict, Any
import logging

Job definitions
def process_bulk_listings(tenant_id: str, listings_data: List[Dict]) -> Dict[str, Any]:
 """Process bulk listing creation for a tenant"""
 results = {
 "tenant_id": tenant_id,
 "total_listings": len(listings_data),
 "successful": 0,
 "failed": 0,
 "errors": []
 }

 etsy_service = EtsyService(tenant_id)

 for listing_data in listings_data:
 try:
 result = await etsy_service.create_listing(listing_data)
 results["successful"] += 1

 # Store success in tenant database
 await store_listing_result(tenant_id, listing_data["id"], result)

 except Exception as e:
 results["failed"] += 1
 results["errors"].append({
 "listing_id": listing_data.get("id"),
 "error": str(e)
 })

 # Log error for tenant
 await log_tenant_error(tenant_id, "listing_creation", str(e))

 return results

def generate_mockup_images(tenant_id: str, design_id: str, mockup_config: Dict) -> Dict[str, Any]:
 """Generate mockup images for a design"""
 import PIL.Image
 import numpy as np
```

```
try:
 # Load design from MinIO
 design_service = DesignService(tenant_id)
 design_data = design_service.get_design(design_id)

 # Load base mockup templates
 mockup_templates = load_mockup_templates(mockup_config["template_type"])

 generated_mockups = []

 for template in mockup_templates:
 # Apply design to template
 mockup_image = apply_design_to_template(
 design_data["image_data"],
 template,
 mockup_config.get("mask_data", {}))
)

 # Save mockup to MinIO
 mockup_path = f"{tenant_id}/mockups/{design_id}_{template['name']}.png"
 mockup_url = design_service.save_mockup(mockup_path, mockup_image)

 generated_mockups.append({
 "template_name": template["name"],
 "mockup_url": mockup_url,
 "size": mockup_image.size
 })

 return {
 "tenant_id": tenant_id,
 "design_id": design_id,
 "generated_mockups": generated_mockups,
 "status": "completed"
 }

except Exception as e:
 logging.error(f"Mockup generation failed for {tenant_id}:{design_id}: {e}")
 return {
 "tenant_id": tenant_id,
 "design_id": design_id,
 "status": "failed",
 "error": str(e)
 }
```

```
def sync_etsy_data(tenant_id: str) -> Dict[str, Any]:
 """Sync Etsy data for tenant (orders, listings, etc.)"""
 etsy_service = EtsyService(tenant_id)

 # Get latest sync timestamp
 last_sync = get_last_sync_timestamp(tenant_id)

 sync_results = {
 "tenant_id": tenant_id,
 "sync_started": datetime.utcnow().isoformat(),
 "orders_synced": 0,
 "listings_synced": 0,
 "errors": []
 }

 try:
 # Sync orders
 new_orders = await etsy_service.get_orders_since(last_sync)
 for order in new_orders:
 await store_tenant_order(tenant_id, order)
 sync_results["orders_synced"] += 1

 # Sync listings
 updated_listings = await etsy_service.get_listings_updates_since(last_sync)
 for listing in updated_listings:
 await store_tenant_listing(tenant_id, listing)
 sync_results["listings_synced"] += 1

 # Update sync timestamp
 await update_sync_timestamp(tenant_id, datetime.utcnow())

 sync_results["status"] = "completed"
 sync_results["sync_completed"] = datetime.utcnow().isoformat()

 except Exception as e:
 sync_results["status"] = "failed"
 sync_results["errors"].append(str(e))
 logging.error(f"Etsy sync failed for {tenant_id}: {e}")

 return sync_results

Worker setup
if __name__ == '__main__':
```

```
redis_conn = redis.from_url(REDIS_URL)

Create queues with priorities
high_priority_queue = Queue('high_priority', connection=redis_conn)
normal_queue = Queue('normal', connection=redis_conn)
low_priority_queue = Queue('low_priority', connection=redis_conn)

Start worker
worker = Worker([high_priority_queue, normal_queue, low_priority_queue],
 connection=redis_conn)
worker.work()
```

## 8.2 Job Management API (Go)

```
go
```

```
// services/jobs/main.go
package main

import (
 "encoding/json"
 "net/http"
 "time"

 "github.com/gin-gonic/gin"
 "github.com/go-redis/redis/v8"
)

type JobService struct {
 redis *redis.Client
 router *gin.Engine
}

type Job struct {
 ID string `json:"id"`
 TenantID string `json:"tenant_id"`
 Type string `json:"type"`
 Status string `json:"status"`
 Payload map[string]interface{} `json:"payload"`
 Result map[string]interface{} `json:"result,omitempty"`
 CreatedAt time.Time `json:"created_at"`
 CompletedAt *time.Time `json:"completed_at,omitempty"`
 Error string `json:"error,omitempty"`
}

func (js *JobService) queueJob(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 var jobRequest struct {
 Type string `json:"type" binding:"required"`
 Payload map[string]interface{} `json:"payload" binding:"required"`
 Priority string `json:"priority" // high, normal, low`
 }

 if err := c.ShouldBindJSON(&jobRequest); err != nil {
 c.JSON(400, gin.H{"error": err.Error()})
 return
 }
}
```

```

// Generate job ID
jobID := generateJobID()

job := Job{
 ID: jobID,
 TenantID: tenantID,
 Type: jobRequest.Type,
 Status: "queued",
 Payload: jobRequest.Payload,
 CreatedAt: time.Now(),
}

// Store job metadata
jobData, _ := json.Marshal(job)
js.redis.Set(c.Request.Context(),
 fmt.Sprintf("job:%s", jobID), jobData, 24*time.Hour)

// Queue job based on type and priority
queueName := js.getQueueForJob(jobRequest.Type, jobRequest.Priority)

jobPayload := map[string]interface{}{
 "job_id": jobID,
 "tenant_id": tenantID,
 "type": jobRequest.Type,
 "payload": jobRequest.Payload,
}

payloadJSON, _ := json.Marshal(jobPayload)

err := js.redis.LPush(c.Request.Context(), queueName, payloadJSON).Err()
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to queue job"})
 return
}

c.JSON(202, gin.H{
 "job_id": jobID,
 "status": "queued",
 "message": fmt.Sprintf("Job queued in %s", queueName)
})

func (js *JobService) getJobStatus(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
}

```

```
jobID := c.Param("job_id")

// Get job data
jobData, err := js.redis.Get(c.Request.Context()),
 fmt.Sprintf("job:%s", jobID)).Result()
if err == redis.Nil {
 c.JSON(404, gin.H{"error": "Job not found"})
 return
}

var job Job
json.Unmarshal([]byte(jobData), &job)

// Verify tenant access
if job.TenantID != tenantID {
 c.JSON(403, gin.H{"error": "Access denied"})
 return
}

c.JSON(200, job)
}

func (js *JobService) listTenantJobs(c *gin.Context) {
 tenantID := c.GetString("tenant_id")

 // Get all job keys for tenant
pattern := fmt.Sprintf("job:*")
keys, err := js.redis.Keys(c.Request.Context(), pattern).Result()
if err != nil {
 c.JSON(500, gin.H{"error": "Failed to fetch jobs"})
 return
}

var jobs []Job
for _, key := range keys {
 jobData, err := js.redis.Get(c.Request.Context(), key).Result()
 if err != nil {
 continue
 }

 var job Job
 if err := json.Unmarshal([]byte(jobData), &job); err != nil {
 continue
 }
}
```

```
// Filter by tenant
if job.TenantID == tenantID {
 jobs = append(jobs, job)
}
}

c.JSON(200, gin.H{
 "jobs": jobs,
 "total": len(jobs)
})
}
```

## 9. Advanced Features Implementation

### 9.1 Real-time Notifications Service (Go + WebSockets)

```
go
```

```
// services/notifications/main.go
package main

import (
 "encoding/json"
 "log"
 "net/http"

 "github.com/gin-gonic/gin"
 "github.com/gorilla/websocket"
)

type NotificationService struct {
 clients map[string]map[*websocket.Conn]bool // tenant_id -> connections
 broadcast chan NotificationMessage
 register chan *Client
 unregister chan *Client
 router *gin.Engine
}

type Client struct {
 TenantID string
 Conn *websocket.Conn
 Send chan NotificationMessage
}

type NotificationMessage struct {
 TenantID string `json:"tenant_id"`
 Type string `json:"type"`
 Title string `json:"title"`
 Message string `json:"message"`
 Data interface{} `json:"data,omitempty"`
 Timestamp int64 `json:"timestamp"`
}

var upgrader = websocket.Upgrader{
 CheckOrigin: func(r *http.Request) bool {
 return true // Configure proper origin checking in production
 },
}

func (ns *NotificationService) handleWebSocket(c *gin.Context) {
 tenantID := c.GetString("tenant_id")
```

```
conn, err := upgrader.Upgrade(c.Writer, c.Request, nil)
if err != nil {
 log.Printf("WebSocket upgrade failed: %v", err)
 return
}

client := &Client{
 TenantID: tenantID,
 Conn: conn,
 Send: make(chan NotificationMessage, 256),
}
}

ns.register <- client

go client.writePump()
go client.readPump(ns.unregister)
}

func (c *Client) writePump() {
 defer c.Conn.Close()

 for {
 select {
 case message, ok := <-c.Send:
 if !ok {
 c.Conn.WriteMessage(websocket.CloseMessage, []byte{})
 return
 }

 if err := c.Conn.WriteJSON(message); err != nil {
 log.Printf("WebSocket write error: %v", err)
 return
 }
 }
 }
}

func (ns *NotificationService) run() {
 for {
 select {
 case client := <-ns.register:
 if ns.clients[client.TenantID] == nil {
 ns.clients[client.TenantID] = make(map[*websocket.Conn]bool)
 }
 }
 }
}
```

```

 }

 ns.clients[client.TenantID][client.Conn] = true
 log.Printf("Client connected for tenant: %s", client.TenantID)

 case client := <-ns.unregister:
 if clients, ok := ns.clients[client.TenantID]; ok {
 if _, ok := clients[client.Conn]; ok {
 delete(clients, client.Conn)
 close(client.Send)
 client.Conn.Close()
 }
 }
 }

 case message := <-ns.broadcast:
 if clients, ok := ns.clients[message.TenantID]; ok {
 for conn := range clients {
 select {
 case conn.WriteJSON(message):
 default:
 close(conn.Send)
 delete(clients, conn)
 conn.Close()
 }
 }
 }
 }

}

// Notification triggers
func (ns *NotificationService) sendJobCompletion(tenantID, jobID, jobType string, result interface{}) {
 message := NotificationMessage{
 TenantID: tenantID,
 Type: "job_completed",
 Title: "Job Completed",
 Message: fmt.Sprintf("%s job completed successfully", jobType),
 Data: map[string]interface{}{"job_id": jobID, "result": result},
 Timestamp: time.Now().Unix(),
 }

 ns.broadcast <- message
}

```

## 9.2 Billing and Subscription Management (Python)

python

```
services/billing/main.py
from fastapi import FastAPI, Depends, HTTPException
from sqlalchemy.orm import Session
from datetime import datetime, timedelta
from typing import Dict, Any
import stripe

app = FastAPI(title="Billing Service")

Subscription tiers configuration
SUBSCRIPTION_TIERS = {
 "starter": {
 "name": "Starter",
 "monthly_price": 29,
 "limits": {
 "max_listings": 100,
 "max_designs": 50,
 "max_mockups_per_month": 500,
 "api_calls_per_hour": 100
 }
 },
 "professional": {
 "name": "Professional",
 "monthly_price": 79,
 "limits": {
 "max_listings": 1000,
 "max_designs": 500,
 "max_mockups_per_month": 5000,
 "api_calls_per_hour": 500
 }
 },
 "enterprise": {
 "name": "Enterprise",
 "monthly_price": 199,
 "limits": {
 "max_listings": -1, # Unlimited
 "max_designs": -1, # Unlimited
 "max_mockups_per_month": -1, # Unlimited
 "api_calls_per_hour": 2000
 }
 }
}
```

```
class BillingService:
 def __init__(self, tenant_id: str, db: Session):
 self.tenant_id = tenant_id
 self.db = db

 @async def get_usage_stats(self) -> Dict[str, Any]:
 """Get current usage statistics for tenant"""

 # Query current month usage
 current_month = datetime.now().replace(day=1, hour=0, minute=0, second=0)

 usage_stats = {
 "current_listings": await self._count_active_listings(),
 "total_designs": await self._count_designs(),
 "mockups_this_month": await self._count_mockups_since(current_month),
 "api_calls_this_hour": await self._count_api_calls_last_hour()
 }

 # Get subscription limits
 subscription = await self._get_tenant_subscription()
 tier_limits = SUBSCRIPTION_TIERS[subscription.tier]["limits"]

 # Calculate usage percentages
 usage_percentages = {}
 for metric, current_value in usage_stats.items():
 limit_key = f"max_{metric}" if metric != "api_calls_this_hour" else "api_calls_per_hour"
 limit = tier_limits.get(limit_key, -1)

 if limit == -1: # Unlimited
 usage_percentages[metric] = 0
 else:
 usage_percentages[metric] = (current_value / limit) * 100

 return {
 "tenant_id": self.tenant_id,
 "subscription_tier": subscription.tier,
 "usage": usage_stats,
 "limits": tier_limits,
 "usage_percentages": usage_percentages,
 "billing_period": {
 "start": subscription.current_period_start,
 "end": subscription.current_period_end
 }
 }
```

```
async def check_usage_limits(self, resource_type: str, requested_amount: int = 1) -> bool:
 """Check if tenant can use more resources"""

 usage_stats = await self.get_usage_stats()
 current_usage = usage_stats["usage"]
 limits = usage_stats["limits"]

 resource_mapping = {
 "listings": ("current_listings", "max_listings"),
 "designs": ("total_designs", "max_designs"),
 "mockups": ("mockups_this_month", "max_mockups_per_month"),
 "api_calls": ("api_calls_this_hour", "api_calls_per_hour")
 }

 if resource_type not in resource_mapping:
 raise ValueError(f"Unknown resource type: {resource_type}")

 usage_key, limit_key = resource_mapping[resource_type]
 current_value = current_usage[usage_key]
 limit = limits[limit_key]

 if limit == -1: # Unlimited
 return True

 return (current_value + requested_amount) <= limit

async def create_stripe_checkout_session(self,
 new_tier: str,
 success_url: str,
 cancel_url: str) -> str:
 """Create Stripe checkout session for subscription upgrade/downgrade"""

 if new_tier not in SUBSCRIPTION_TIERS:
 raise HTTPException(400, f"Invalid subscription tier: {new_tier}")

 tenant = await self._get_tenant()
 price_amount = SUBSCRIPTION_TIERS[new_tier]["monthly_price"] * 100 # Stripe uses cents

 checkout_session = stripe.checkout.Session.create(
 customer_email=tenant.primary_email,
 payment_method_types=['card'],
 line_items=[{
 'price_data': {
```

```

'currency': 'usd',
'unit_amount': price_amount,
'product_data': {
 'name': f'Etsy Seller Automater - {SUBSCRIPTION_TIERS[new_tier]["name"]}',
 'description': f'Monthly subscription to {new_tier} tier'
},
'recurring': {
 'interval': 'month'
},
'quantity': 1,
}],
mode='subscription',
success_url=success_url + '?session_id={CHECKOUT_SESSION_ID}',
cancel_url=cancel_url,
metadata={
 'tenant_id': self.tenant_id,
 'subscription_tier': new_tier
}
)

return checkout_session.url

@app.get("/billing/usage")
async def get_tenant_usage(tenant_id: str = Depends(get_tenant_id)):
 """Get detailed usage statistics for tenant"""
 billing_service = BillingService(tenant_id, get_db())
 return await billing_service.get_usage_stats()

@app.post("/billing/upgrade")
async def create_upgrade_session(
 new_tier: str,
 tenant_id: str = Depends(get_tenant_id)
):
 """Create Stripe checkout session for subscription upgrade"""
 billing_service = BillingService(tenant_id, get_db())

 checkout_url = await billing_service.create_stripe_checkout_session(
 new_tier=new_tier,
 success_url=f"https://{{tenant_id}}.yourdomain.com/billing/success",
 cancel_url=f"https://{{tenant_id}}.yourdomain.com/billing/cancel"
)

 return {"checkout_url": checkout_url}

```

```
@app.post("/billing/webhook")
async def stripe_webhook(request: Request):
 """Handle Stripe webhooks for subscription events"""
 payload = await request.body()
 sig_header = request.headers.get('Stripe-Signature')

 try:
 event = stripe.Webhook.construct_event(
 payload, sig_header, STRIPE_WEBHOOK_SECRET
)
 except ValueError:
 raise HTTPException(400, "Invalid payload")
 except stripe.error.SignatureVerificationError:
 raise HTTPException(400, "Invalid signature")

 if event['type'] == 'checkout.session.completed':
 session = event['data']['object']
 tenant_id = session['metadata']['tenant_id']
 new_tier = session['metadata']['subscription_tier']

 # Update tenant subscription
 await update_tenant_subscription(tenant_id, new_tier, session)

 elif event['type'] == 'invoice.payment_failed':
 # Handle failed payment
 invoice = event['data']['object']
 await handle_payment_failure(invoice)

 return {"status": "success"}
```

## 10. Performance Optimization Strategies

### 10.1 Database Query Optimization

python

```

services/common/database_optimizations.py
from sqlalchemy import create_engine, event
from sqlalchemy.orm import sessionmaker
from sqlalchemy.pool import QueuePool
import logging

Connection pooling configuration for QNAP constraints
DATABASE_CONFIG = {
 "pool_size": 5, # Limit connections due to QNAP resources
 "max_overflow": 10, # Maximum additional connections
 "pool_pre_ping": True, # Verify connections before use
 "pool_recycle": 3600, # Recycle connections every hour
}

engine = create_engine(
 DATABASE_URL,
 **DATABASE_CONFIG,
 echo=False # Set to True for query debugging
)

Query performance monitoring
@event.listens_for(engine, "before_cursor_execute")
def receive_before_cursor_execute(conn, cursor, statement, parameters, context, executemany):
 context._query_start_time = time.time()

@event.listens_for(engine, "after_cursor_execute")
def receive_after_cursor_execute(conn, cursor, statement, parameters, context, executemany):
 total = time.time() - context._query_start_time
 if total > 1.0: # Log slow queries
 logging.warning(f"Slow query ({total:.2f}s): {statement[:200]}...")

Optimized tenant data access patterns
class TenantDataAccess:
 def __init__(self, tenant_id: str, db: Session):
 self.tenant_id = tenant_id
 self.db = db

 def get_dashboard_data(self) -> Dict[str, Any]:
 """Optimized dashboard data retrieval with single query"""

 # Single query to get all dashboard metrics
 result = self.db.execute(text("""
 WITH tenant_stats AS (

```

```

SELECT
 COUNT(DISTINCT o.id) as total_orders,
 COALESCE(SUM(o.total_amount), 0) as total_revenue,
 COUNT(DISTINCT p.id) as total_products,
 COUNT(DISTINCT d.id) as total_designs
FROM {schema}.orders o
FULL OUTER JOIN {schema}.products p ON true
FULL OUTER JOIN {schema}.designs d ON true
WHERE o.created_at >= CURRENT_DATE - INTERVAL '30 days'
),
recent_activity AS (
SELECT
 'order' as activity_type,
 id::text as activity_id,
 created_at,
 total_amount::text as activity_data
FROM {schema}.orders
WHERE created_at >= CURRENT_DATE - INTERVAL '7 days'

UNION ALL

SELECT
 'design' as activity_type,
 id::text as activity_id,
 created_at,
 name as activity_data
FROM {schema}.designs
WHERE created_at >= CURRENT_DATE - INTERVAL '7 days'
ORDER BY created_at DESC
LIMIT 10
)
SELECT
 json_build_object(
 'stats', row_to_json(ts),
 'recent_activity', json_agg(ra)
) as dashboard_data
FROM tenant_stats ts, recent_activity ra
GROUP BY ts.total_orders, ts.total_revenue, ts.total_products, ts.total_designs
""".format(schema=f"tenant_{self.tenant_id}"))

```

**return result.fetchone()[0]**

## 10.2 Caching Strategy (Redis)

python

```
services/common/cache.py
import redis
import json
import hashlib
from functools import wraps
from typing import Any, Optional, Callable
import pickle

redis_client = redis.Redis.from_url(REDIS_URL, decode_responses=False)

class TenantCache:
 def __init__(self, tenant_id: str):
 self.tenant_id = tenant_id
 self.key_prefix = f"tenant:{tenant_id}"

 def cache_key(self, key: str) -> str:
 """Generate tenant-scoped cache key"""
 return f"{self.key_prefix}:{key}"

 def get(self, key: str) -> Optional[Any]:
 """Get cached value"""
 cache_key = self.cache_key(key)
 data = redis_client.get(cache_key)
 if data:
 return pickle.loads(data)
 return None

 def set(self, key: str, value: Any, expire: int = 3600) -> None:
 """Set cached value with expiration"""
 cache_key = self.cache_key(key)
 redis_client.setex(cache_key, expire, pickle.dumps(value))

 def delete(self, key: str) -> None:
 """Delete cached value"""
 cache_key = self.cache_key(key)
 redis_client.delete(cache_key)

 def invalidate_pattern(self, pattern: str) -> None:
 """Invalidate all keys matching pattern"""
 pattern_key = self.cache_key(pattern)
 keys = redis_client.keys(pattern_key)
 if keys:
 redis_client.delete(*keys)
```

```
def cache_tenant_data(expire: int = 3600, key_suffix: str = None):
 """Decorator for caching tenant-specific data"""
 def decorator(func: Callable) -> Callable:
 @wraps(func)
 async def wrapper(tenant_id: str, *args, **kwargs):
 cache = TenantCache(tenant_id)

 # Generate cache key from function name and parameters
 if key_suffix:
 cache_key = f"{func.__name__}:{key_suffix}"
 else:
 # Create hash of parameters for unique key
 params_hash = hashlib.md5(
 json.dumps({"args": args, "kwargs": kwargs}, sort_keys=True).encode()
).hexdigest()[:8]
 cache_key = f"{func.__name__}:{params_hash}"

 # Try to get from cache
 cached_result = cache.get(cache_key)
 if cached_result is not None:
 return cached_result

 # Execute function and cache result
 result = await func(tenant_id, *args, **kwargs)
 cache.set(cache_key, result, expire)

 return result
 return wrapper
 return decorator

Usage examples
@cache_tenant_data(expire=1800, key_suffix="analytics")
async def get_tenant_analytics(tenant_id: str, date_range: str) -> Dict[str, Any]:
 """Cached analytics data"""
 # Expensive analytics computation here
 pass

@cache_tenant_data(expire=300) # 5 minutes for frequently changing data
async def get_etsy_listings(tenant_id: str, page: int = 1) -> Dict[str, Any]:
 """Cached Etsy listings with pagination"""
 # API call to Etsy
 pass
```

## 11. Deployment and Operations

### 11.1 Production Deployment Script for QNAP

```
bash
```

```
#!/bin/bash
deploy-production.sh

set -e

QNAP_HOST="your-qnap-ip"
QNAP_USER="admin"
PROJECT_PATH="/share/Container/etsy-saas-prod"
BACKUP_PATH="/share/Container/backups"

echo "🚀 Starting production deployment to QNAP NAS..."

Create deployment directory
ssh $QNAP_USER@$QNAP_HOST "mkdir -p $PROJECT_PATH"

Backup current deployment if exists
ssh $QNAP_USER@$QNAP_HOST "
if [-d $PROJECT_PATH/current]; then
 echo '📦 Creating backup of current deployment...'
 mkdir -p $BACKUP_PATH
 tar -czf $BACKUP_PATH/backup-$(date +%Y%m%d-%H%M%S).tar.gz -C $PROJECT_PATH current
fi
"

Upload new deployment files
echo "📦 Uploading deployment files..."
rsync -avz --exclude='\.git' --exclude='node_modules' --exclude='__pycache__' \
./ $QNAP_USER@$QNAP_HOST:$PROJECT_PATH/new/

Switch to new deployment
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH
if [-d current]; then
 mv current previous
fi
mv new current
"

Start services
echo "⌚ Starting production services..."
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH/current
docker-compose -f docker-compose.prod.yml down || true
```

```
docker-compose -f docker-compose.prod.yml pull
docker-compose -f docker-compose.prod.yml up -d
""

Health check
echo "🏥 Running health checks..."
sleep 30
ssh $QNAP_USER@$QNAP_HOST "
cd $PROJECT_PATH/current
./scripts/health-check.sh
""

Cleanup old images
ssh $QNAP_USER@$QNAP_HOST "
echo '🧹 Cleaning up old Docker images...'
docker system prune -f
docker image prune -f
""

echo "✅ Production deployment completed successfully!"
echo "🌐 Services available at: https://your-domain.com"
echo "💻 Monitor logs with: ssh $QNAP_USER@$QNAP_HOST 'cd $PROJECT_PATH/current && docker-compos
```